

Assignment #1: The Big Dot

The dot product of two vectors $a = (a_0, a_1, \dots, a_{n-1})$ and $b = (b_0, b_1, \dots, b_{n-1})$, written $a \cdot b$, is simply the sum of the component-by-component products:

$$a \cdot b = \sum_{i=0}^{n-1} a_i \times b_i$$

Dot products are used extensively in computing and have a wide range of applications. For instance, in 3D graphics ($n = 3$), we often make use of the fact that $a \cdot b = |a||b|\cos\theta$, where $| \ |$ denotes vector length and θ is the angle between the two vectors.

In this assignment, you are expected to:

1. Write CUDA code to compute in parallel the dot product of two (possibly large $N = 100,000$, or $N = 1024 \times 1024$) random single precision floating point vectors;
2. Write two functions to compute the results on the CPU and GPU, and compare the two results to check for correctness ($1.0e-6$);
 - `float *CPU_big_dot(float *A, float *B, int N);`
 - `float *GPU_big_dot(float *A, float *B, int N);`
3. Print performance statistics with timer function;
 - CPU: `Tcpu = Total computation time for CPU_big_dot();`
 - GPU: `Tgpu = Total computation time for GPU_big_dot();`
 - Memory allocation and data transfer from CPU to GPU time
 - Kernel execution time
 - Data transfer from GPU to CPU time
 - `Speedup = CPU/GPU`
4. Analyze the performance results in a few sentences.
 - Which one runs faster?

- What's the reason for that? Problem size, overhead, etc.

Timer functions

```
#include <sys/time.h>
```

```
long long start_timer() {  
    struct timeval tv;  
    gettimeofday(&tv, NULL);  
    return tv.tv_sec * 1000000 + tv.tv_usec;  
}
```

```
long long stop_timer(long long start_time, char *name) {  
    struct timeval tv;  
    gettimeofday(&tv, NULL);  
    long long end_time = tv.tv_sec * 1000000 + tv.tv_usec;  
    Printf("%s: %.5f sec\n", name, ((float) (end_time - start_time)) /  
        (1000 * 1000));  
    return end_time - start_time;  
}
```

Due by 5/29/2020 11:59pm