

CPSC/ECE 4780/6780

General-Purpose Computation on Graphical Processing Units (GPGPU)

Lecture 3: Introduction to CUDA

Recaps from Last Lecture

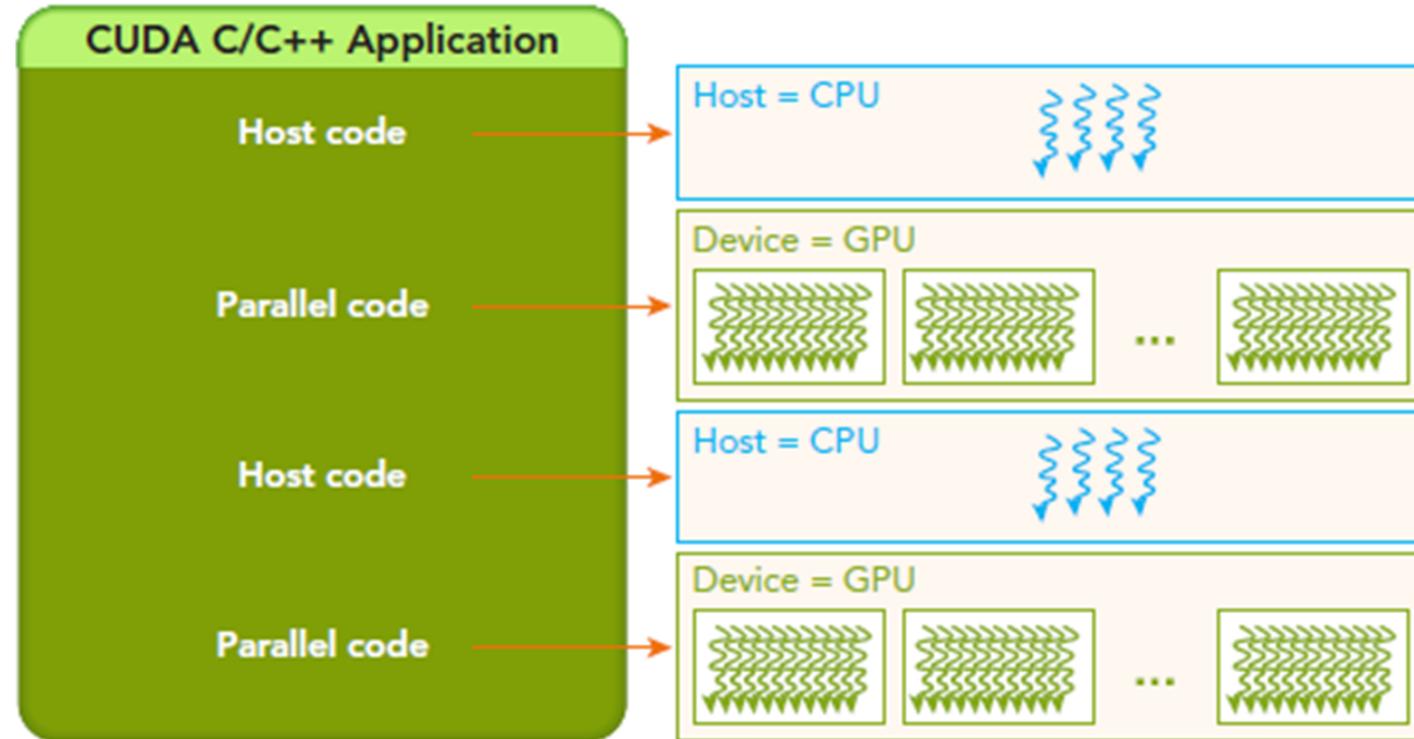
- What is GPU?
- History of GPUs
- Architecture of GPU
- CPU/GPU comparisons
- Why should we use GPUs?
- CPU+GPU acceleration
- GPGPU programming

What is CUDA?

- CUDA – “Compute Unified Device Architecture”
- General-purpose parallel computing platform and programming model
- Created by NVIDIA first in 2007
- Written mostly like C

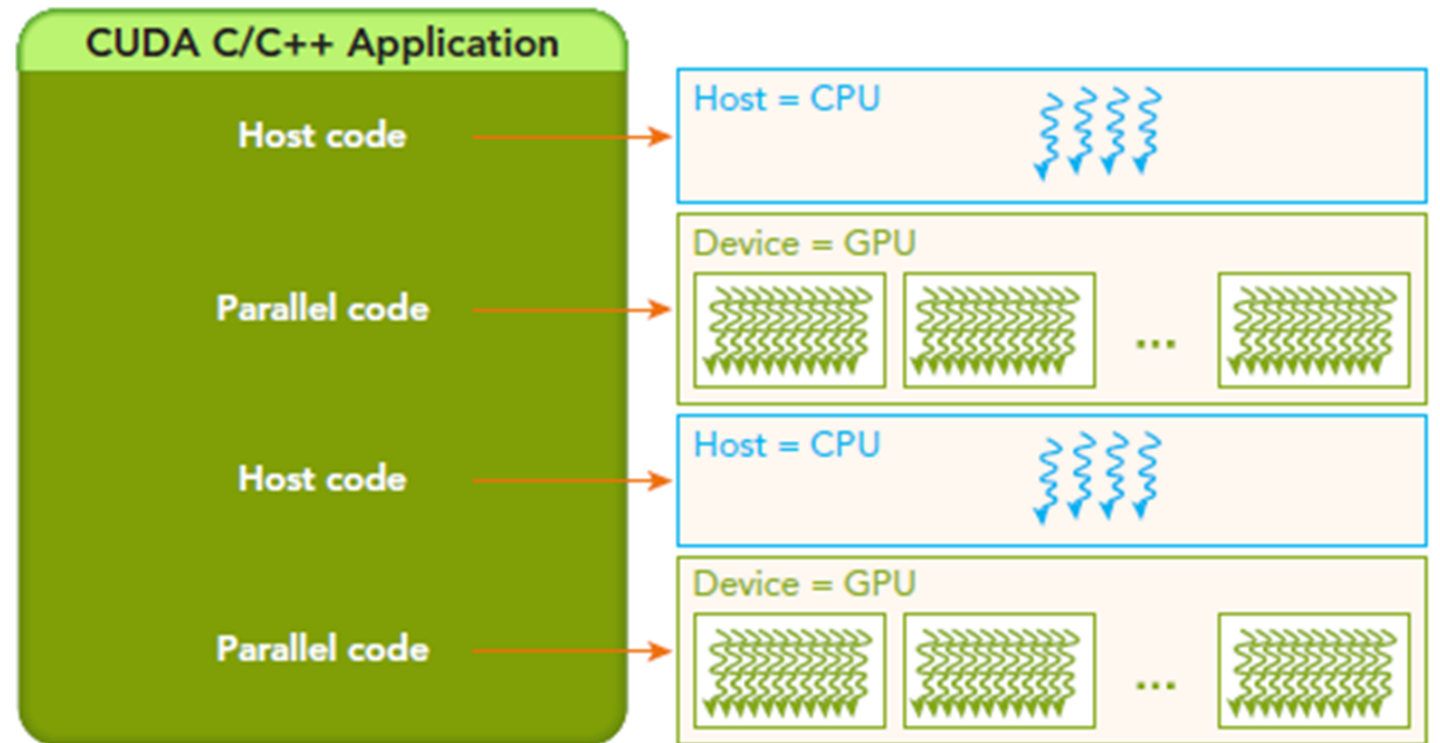
CUDA Programming Structure

- Integrated host + device application C program
 - Host – CPU and its memory
 - Serial or modestly parallel parts
 - Written in ANSI C
 - Device – GPU and its memory
 - Highly parallel parts
 - Written in CUDA C
 - “Kernel”



Processing Flow of a CUDA Program

- Copy input data from CPU memory to GPU memory
- Invoke kernels to operate on the data stored in GPU memory
- Copy data back from GPU memory to CPU memory



Memory Management and Data Transfer

- Host and device memory are separate entities
 - Host pointers point to CPU memory
 - May be passed to/from device code
 - May not be dereferenced in device code
 - Device pointers point to GPU memory
 - May be passed to/from host code
 - May not be dereferenced in host code

STANDARD C FUNCTIONS	CUDA C FUNCTIONS
<code>malloc</code>	<code>cudaMalloc</code>
<code>memcpy</code>	<code>cudaMemcpy</code>
<code>memset</code>	<code>cudaMemset</code>
<code>free</code>	<code>cudaFree</code>

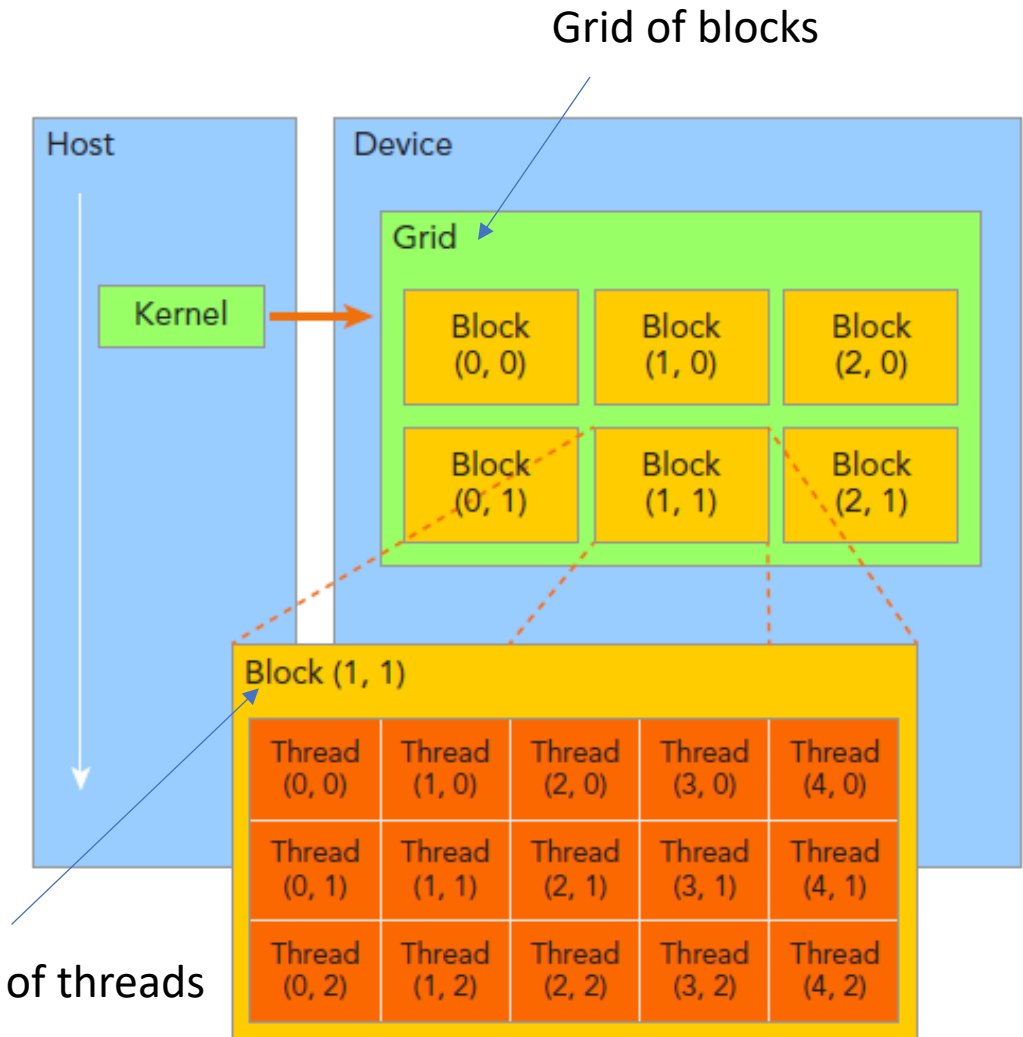
Host and device memory functions

CUDA Function Declaration

- CUDA extensions to C functional declaration
 - `__global__`: indicates a CUDA kernel function
 - executed on the device
 - Only callable from the host
 - Must have a void return type
 - `__device__`: indicates a CUDA device function
 - Executed on the device
 - Only callable from the device
 - `__host__`: indicates a CUDA host function
 - Executed on the host
 - Only callable from the host

Organizing Threads

- Two-level thread hierarchy
 - Grids of blocks
 - Blocks of threads
- All threads in a grid share the same global memory space
- A thread block is a group of threads that can cooperate with each other by:
 - Block-local synchronization
 - Block-local shared memory
- Threads coordinates:
 - blockIdx (block index within a grid)
 - threadIdx (thread index within a block)
 - Type: uint3 (.x, .y, .z)

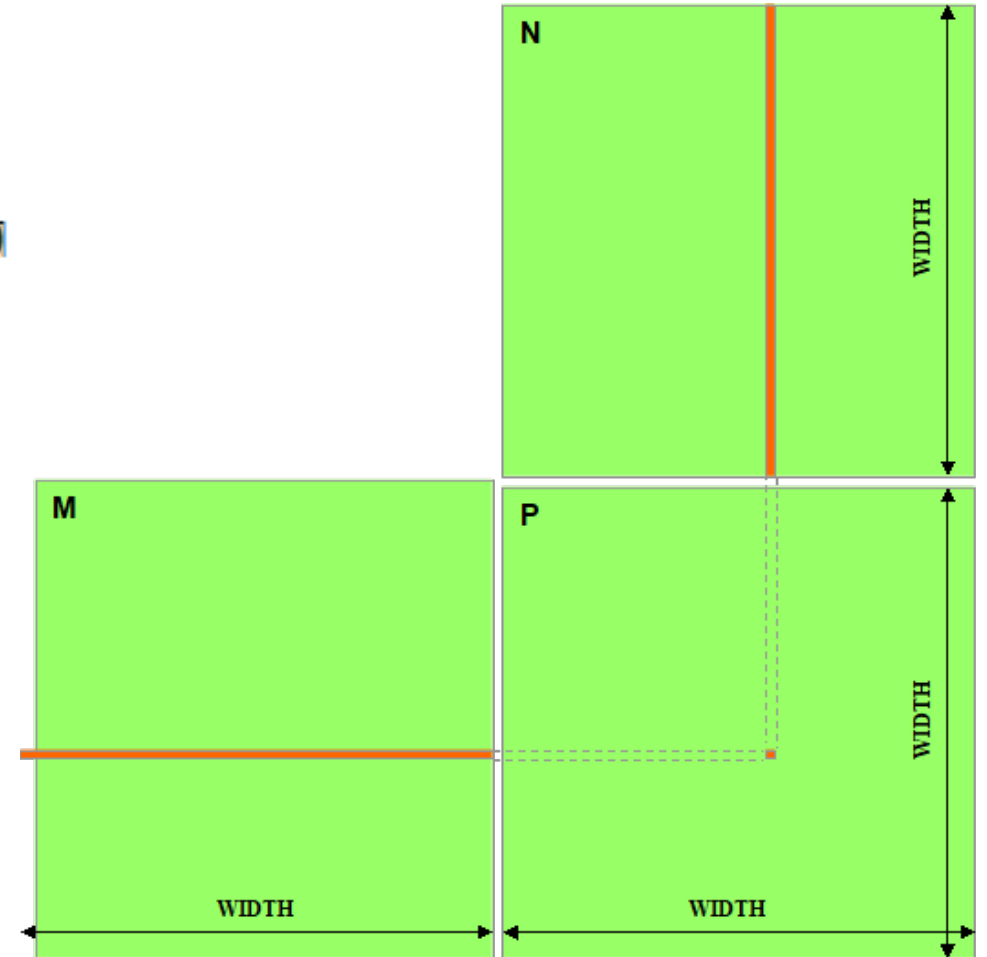


A thread hierarchy structure with a 2D grid containing 2D blocks

Matrix Multiplication on CPU

- $M * N \Rightarrow P$

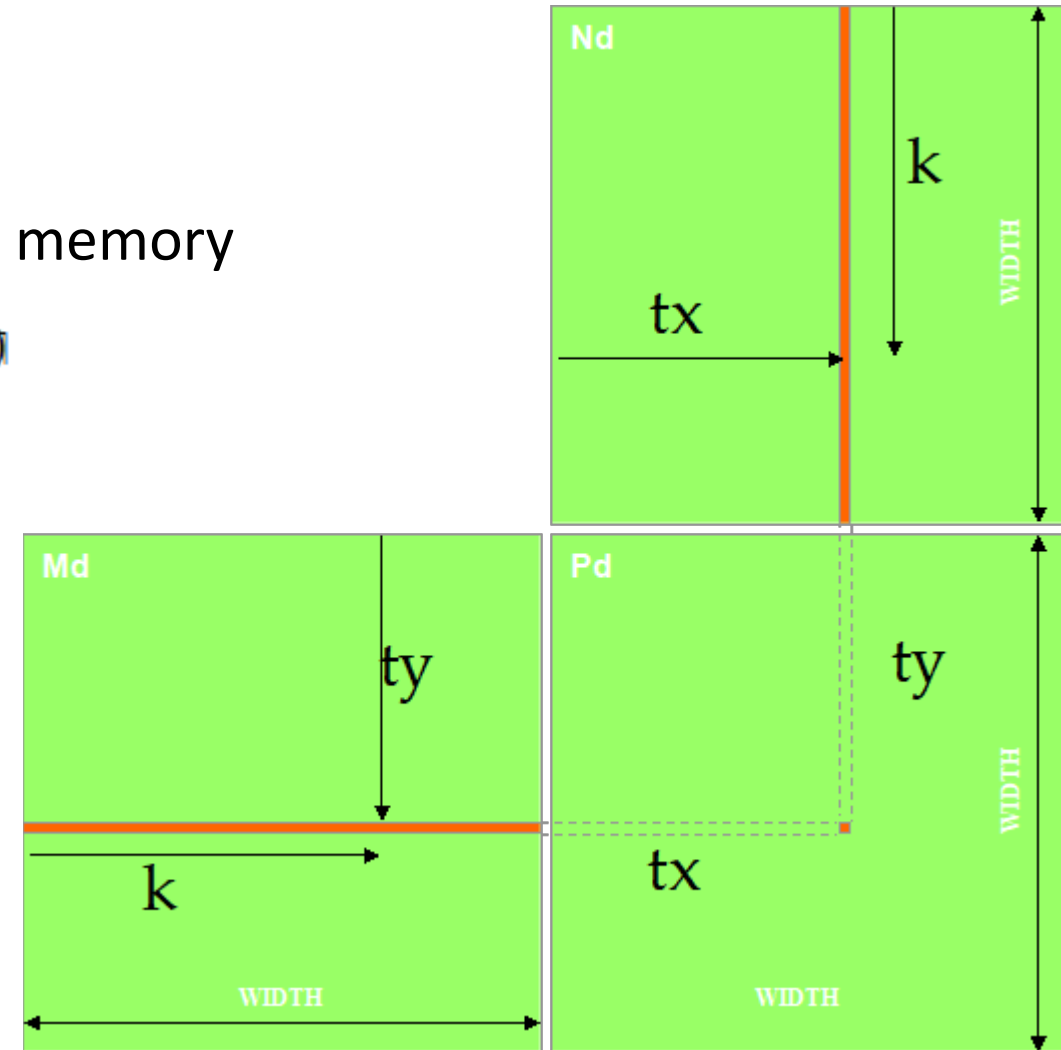
```
void MatrixMulOnHost(float* M, float* N, float* P, int Width)
{
    for (int i = 0; i < Width; ++i)
        for (int j = 0; j < Width; ++j) {
            double sum = 0;
            for (int k = 0; k < Width; ++k) {
                double a = M[i * width + k];
                double b = N[k * width + j];
                sum += a * b;
            }
            P[i * Width + j] = sum;
        }
}
```



Matrix Multiplication on GPU

- One thread calculates one element of P
- M and N are loaded width times from global memory

```
__global__ void MatrixMulKernel(float* Md, float* Nd, float* Pd, int Width)
{
    int tx = threadIdx.x;
    int ty = threadIdx.y;
    float Pvalue = 0;
    for (int k = 0; k < Width; ++k) {
        float Melement = Md[threadIdx.y*Width+k];
        float Nelement = Nd[k*Width+threadIdx.x];
        Pvalue += Melement * Nelement;
    }
    Pd[threadIdx.y*Width+threadIdx.x] = Pvalue;
}
```



CUDA Device Properties

- Get the count of CUDA devices
 - `int count;`
 - `cudaGetDeviceCount(&count);`
- Query relevant information of a device
 - `cudaDeviceProp prop;`
 - `cudaGetDeviceProperties(&prop, i);`
- Set device property and choose a proper device
 - `int dev;`
 - `cudaDeviceProp prop;`
 - `prop.major = 1;`
 - `prop.minor = 3`
 - `cudaChooseDevice(&dev, &prop);`
 - `cudaSetDevice(dev);`

Coding Examples

- Coding

- First CUDA program: Hello World
- Add two numbers
- Add two vectors
 - By blockIdx
 - By threadIdx
 - Combined
- Query device property

- Compilation: use nvcc

- Makefile

Palmetto is comprised of 2021 compute nodes (totalling 23072 CPU cores), and features:

- 2021 compute nodes, totaling 23072 cores
- 595 nodes equipped with 2x NVIDIA Tesla GPUs (2 per node); 103 nodes each have 2x NVIDIA Tesla V100 GPUs (2 per node)

Login with command:

- `ssh username@login.palmetto.clemson.edu`

https://www.palmetto.clemson.edu/palmetto/userguide_palmetto_overview.html

Request a specific GPU (m2075, m2070q, k20, k40, p100, or v100) on Palmetto with command:

- `qsub -l -l select=1:ncpus=1:ngpus=1:gpu_model=k20:mem=2gb,walltime=2:00:00`

https://www.palmetto.clemson.edu/palmetto/userguide_howto_use_gpus.html

Load CUDA module:

- `module load cuda-toolkit`

GPU device query:

- `/software/cuda-toolkit/8.0.44/samples/1_Uutilities/deviceQuery/deviceQuery`

Compile .cu code with “nvcc”, e.g.,

- `nvcc helloWorld.cu`
- `nvcc -o helloWorld helloWorld.cu`

Assignment #1: The Big Dot

- The dot product of two vectors $a = (a_0, a_1, \dots, a_{n-1})$ and $b = (b_0, b_1, \dots, b_{n-1})$, written $a \cdot b$, is simply the sum of the component-by-component products:

$$a \cdot b = \sum_{i=0}^{n-1} a_i \times b_i$$

Dot products are used extensively in computing and have a wide range of applications. For instance, in 3D graphics ($n = 3$), we often make use of the fact that $a \cdot b = |a||b|\cos\theta$, where $| \ |$ denotes vector length and θ is the angle between the two vectors.

Assignment #1: The Big Dot

- Write CUDA code to compute in parallel the dot product of two (possibly large $N = 100,000$, or $N = 1024 \times 1024$) random single precision floating point vectors;
- Write two functions to compute the results on the CPU and GPU, and compare the two results to check for correctness ($1.0e-6$);
 - `float *CPU_big_dot(float *A, float *B, int N);`
 - `float *GPU_big_dot(float *A, float *B, int N);`
- Print performance statistics with timer function;
 - CPU: T_{cpu} = Total computation time for `CPU_big_dot()`;
 - GPU: T_{gpu} = Total computation time for `GPU_big_dot()`;
 - Memory allocation and data transfer from CPU to GPU time
 - Kernel execution time
 - Data transfer from GPU to CPU time
 - $Speedup = GPU/CPU$
- Analyze the performance results in a few sentences.
 - Which one runs faster?
 - What's the reason for that? Problem size, overhead, etc.

Assignment #1: The Big Dot

- Timer functions

- `#include <sys/time.h>`
- `long long start_timer() {
 struct timeval tv;
 gettimeofday(&tv, NULL);
 return tv.tv_sec * 1000000 + tv.tv_usec;
}`
- `long long stop_timer(long long start_time, char *name) {
 struct timeval tv;
 gettimeofday(&tv, NULL);
 long long end_time = tv.tv_sec * 1000000 + tv.tv_usec;
 Printf("%s: %.5f sec\n", name, ((float) (end_time - start_time)) / (1000 * 1000));
 return end_time - start_time;
}`