

CPSC/ECE 4780/6780

General-Purpose Computation on Graphical Processing Units (GPGPU)

Lecture 9: Streams

Recap of Last Lecture

- What are race conditions?
- What is atomic operation?
- What kind of atomic operations do we have?
- What is atomic lock?
- How to implement lock function?

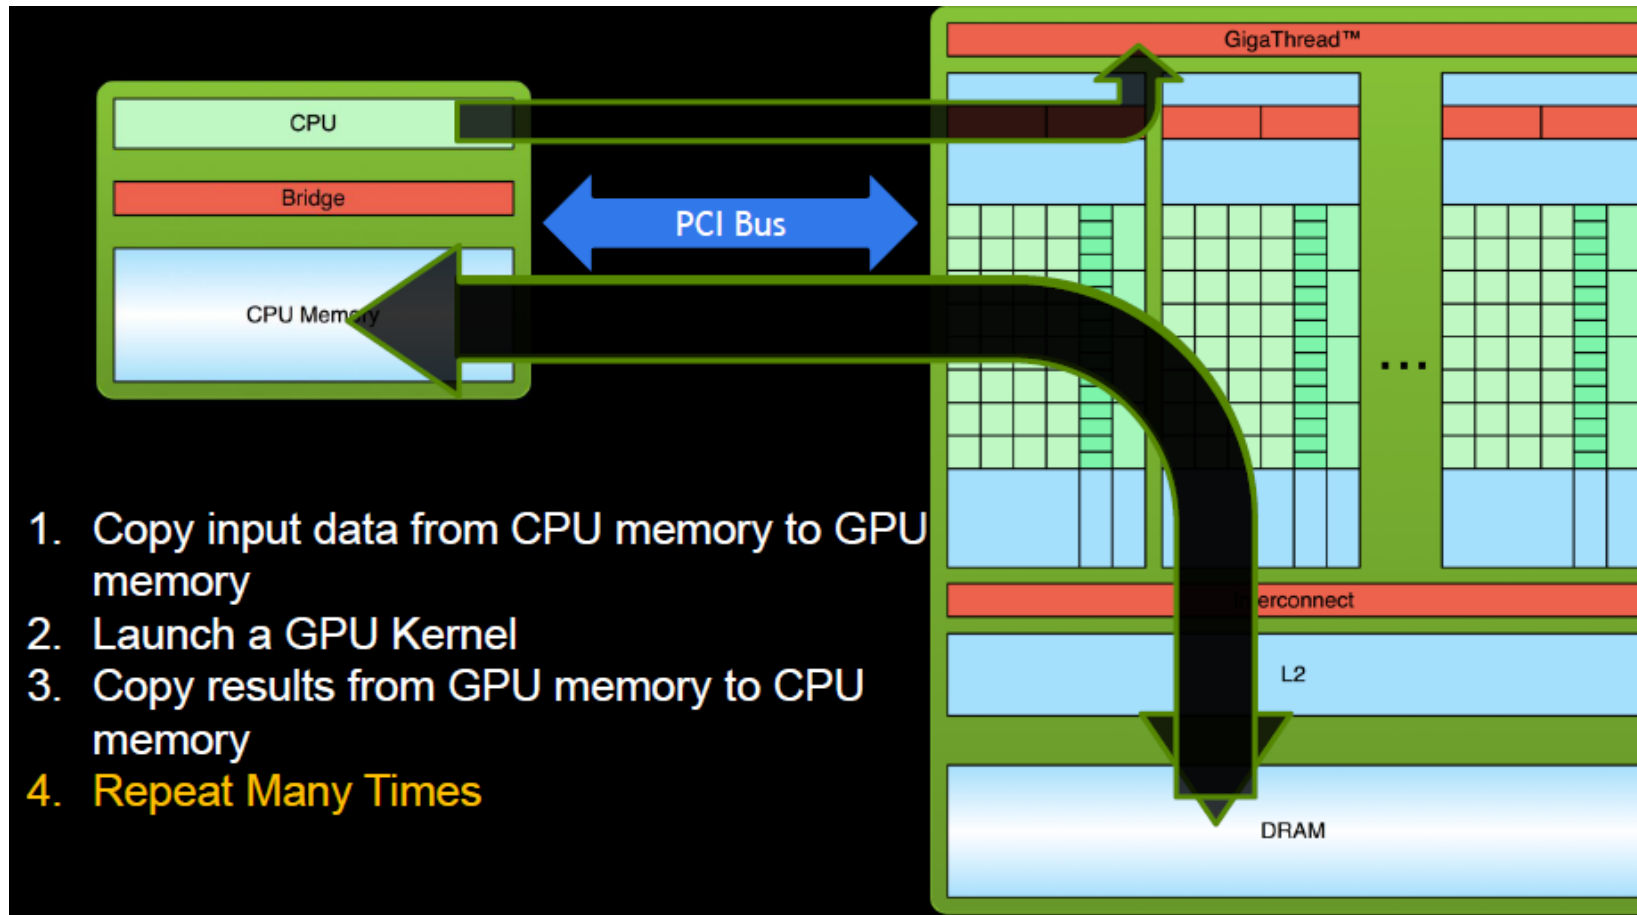
Concurrency

- The ability to perform multiple CUDA operations simultaneously
- Two levels of concurrency in CUDA C programming
 - Kernel level concurrency:
 - A single task, or kernel, is executed in parallel by many threads on the GPU
 - Grid level concurrency:
 - Multiple kernel launches are executed simultaneously on a single device

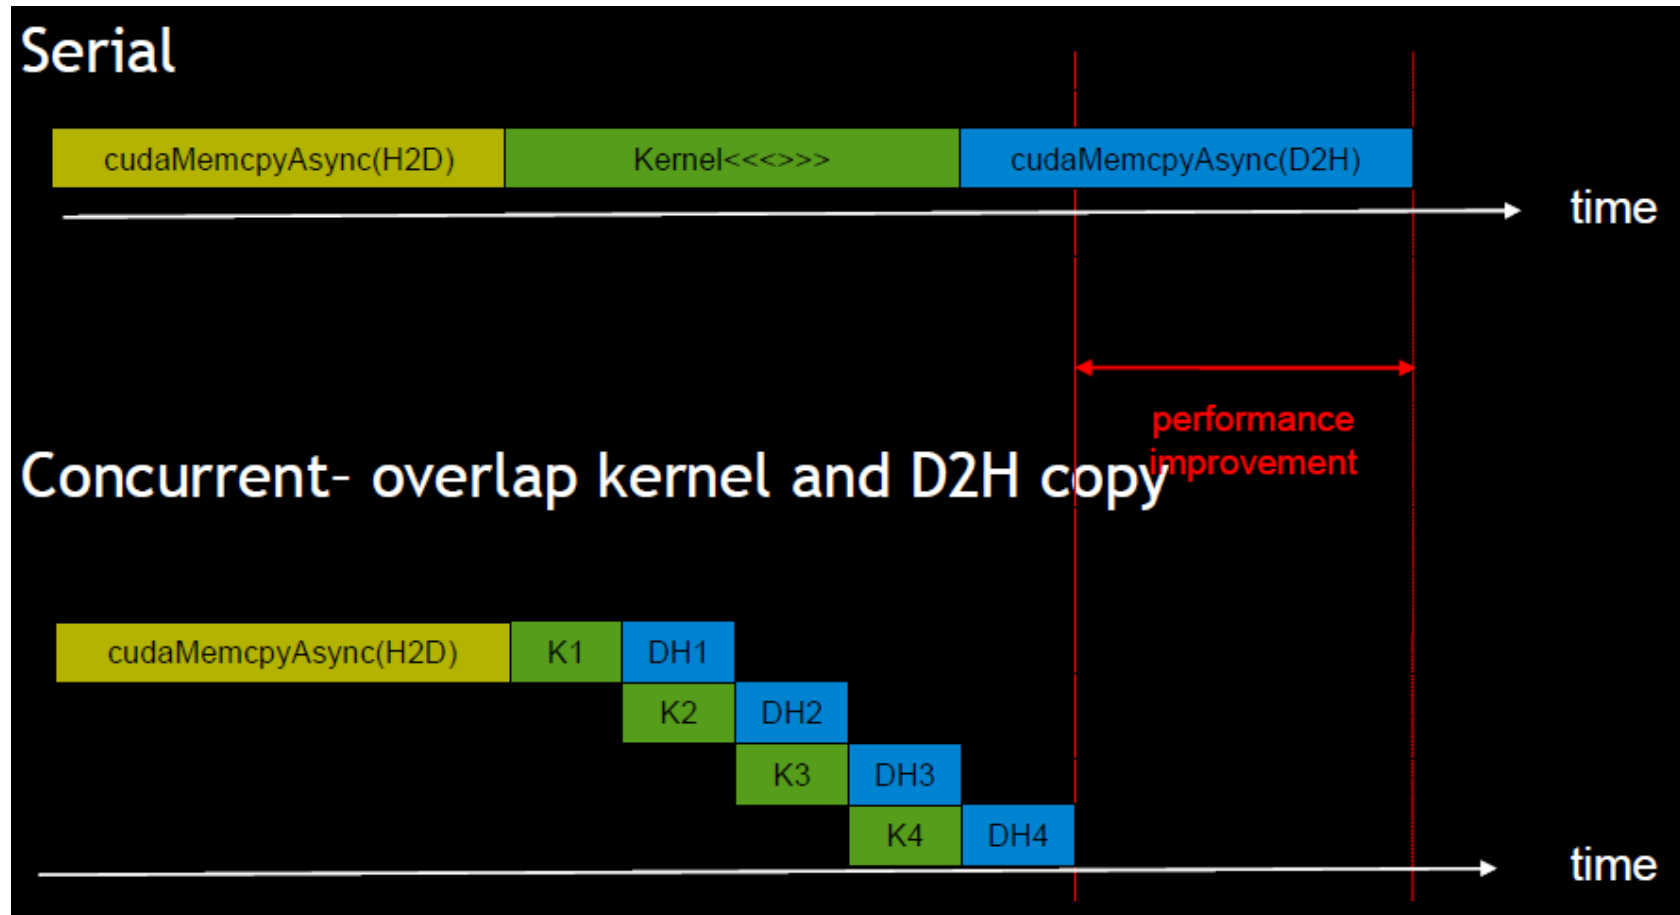
CUDA Streams

- A stream is a queue of device work
 - The host places work in the queue and continues on immediately
 - Device schedules work from streams when resources are free
- CUDA operations are encapsulated in a stream
 - E.g., host-device data transfer, kernel launches, and etc
- Operations within the same stream are ordered (FIFO) and cannot overlap
- Operations in different streams are unordered and can overlap

Serial Processing Flow without Streams



Concurrent Processing Flow with Streams



Streams and Concurrency

- All CUDA operations (both kernels and data transfers) either explicitly or implicitly run in a stream
 - Implicitly declared stream (**NULL stream**): default stream
 - Explicitly declared stream (**non-NULL stream**)
- **Asynchronous, stream-based** kernel launches and data transfers enable four types of concurrency:
 - Overlapped host computation and device computation
 - Overlapped host computation and host-device data transfer
 - Overlapped host-device data transfer and device computation
 - Concurrent device computation

NULL Stream

- Consider the following code using NULL stream:

```
cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice);  
kernel<<<grid, block>>>(d_a);  
cudaMemcpy(h_a, d_a, size, cudaMemcpyDeviceToHost);
```

- From the device perspective
 - All three operations are executed in order on the stream
 - No awareness of any other host operations being performed
- From the host perspective
 - Each data transfer is synchronous
 - Kernel launch is asynchronous => overlap device and host computation

```
cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice);  
Kernel<<<grid, block>>>(d_a);  
anyCPUfunction();  
cudaMemcpy(h_a, d_a, size, cudaMemcpyDeviceToHost);
```


Non-NULL Stream

- Non-NULL streams in CUDA are declared, created, and destroyed in host code as follows:
 - **`cudaStream_t stream;`** // Declare a stream handle
 - **`cudaStreamCreate(&stream);`** // Allocate a stream
 - **`cudaStreamDestroy(stream);`** // Deallocate a stream
- To issue data transfer to non-NULL stream
 - **`cudaMemcpyAsync(d_a, h_a, size, cudaMemcpyHostToDevice, stream)`**
 - **`cudaMemcpyAsync(h_a, d_a, size, cudaMemcpyDeviceToHost, stream)`**
- To launch a kernel to non-NULL stream
 - **`Kernel<<<grid, block, sharedMemSize, stream>>>(d_a);`**

Synchronize and Query in Non-NULL Stream

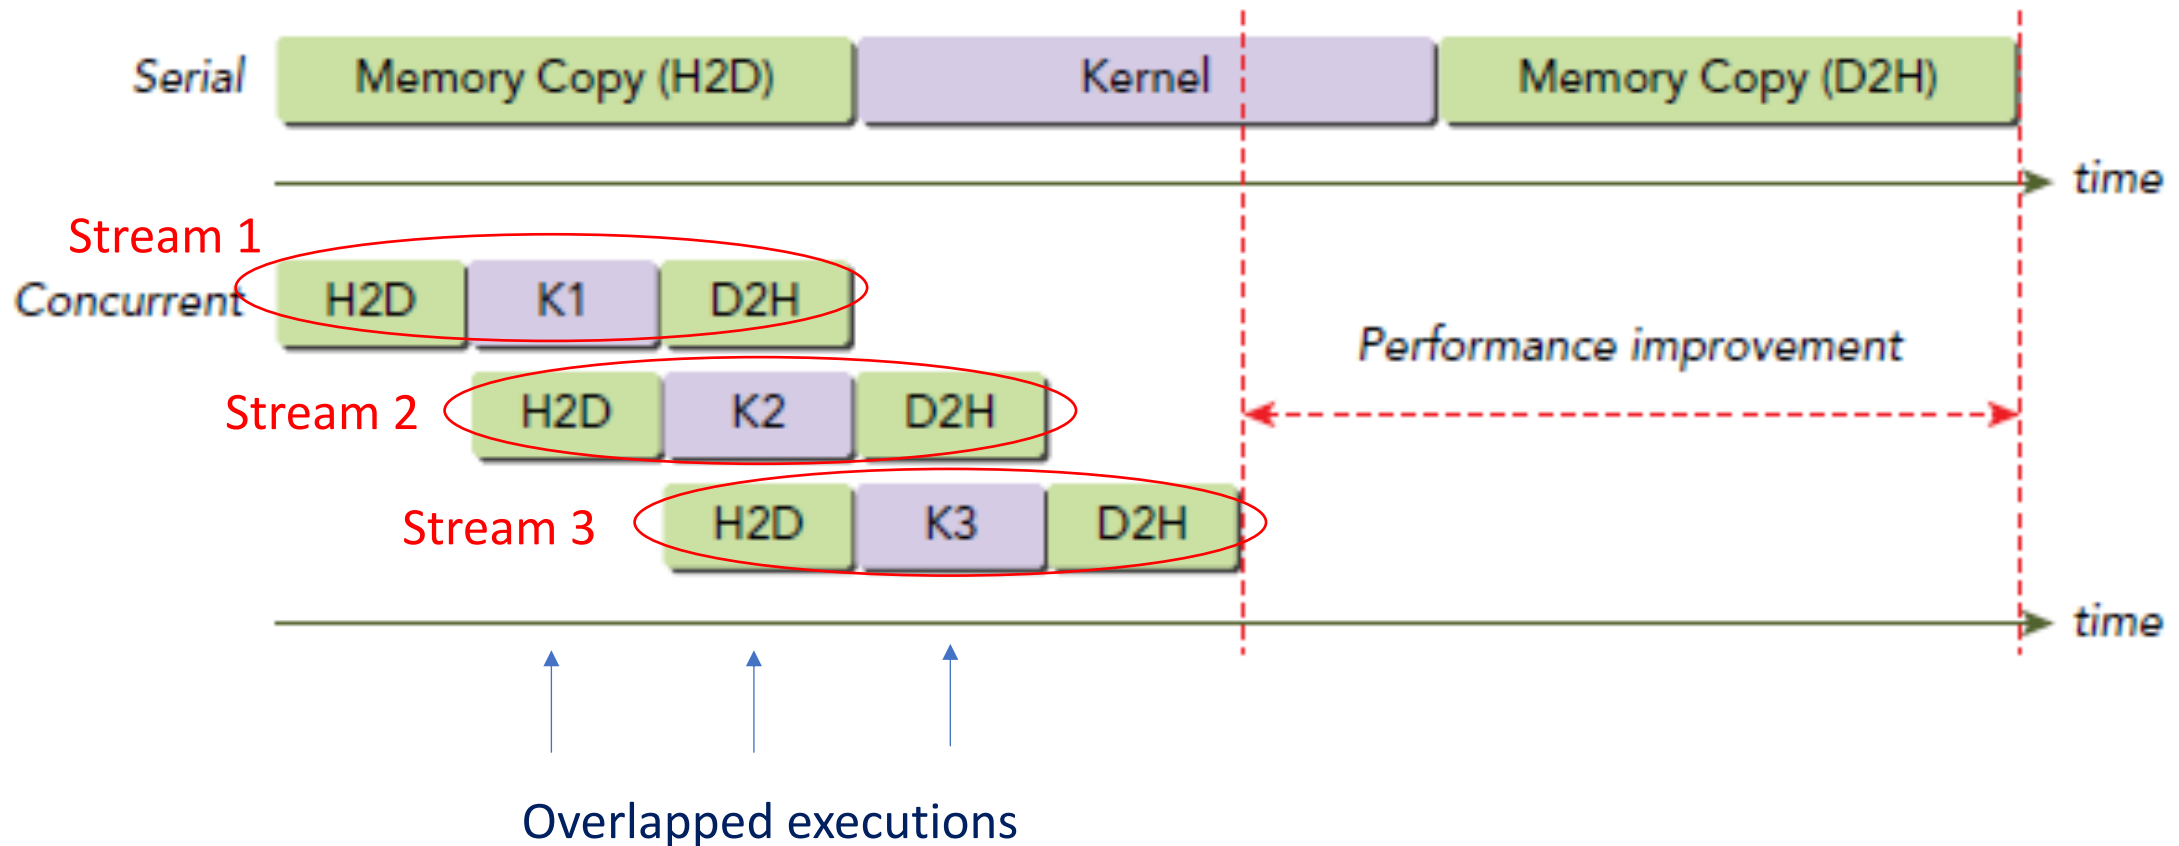
- All operations in non-NULL streams are non-blocking with respect to the host
- Sometimes you need to synchronize the operations with the host
 - **`cudaStreamSynchronize(stream);`** => blocks host
- Sometimes you want to check if all operations in a stream have completed, but does not want to block the host if they have not completed
 - **`cudaStreamQuery(stream);`** => does not block host
 - Returns `cudaSuccess` if all operations are complete
 - Returns `cudaErrorNotReady` otherwise

A Common Pattern for Dispatching CUDA Operations to Multiple Streams

```
for (int i = 0; i < nStreams; i++) {
    int offset = i * bytesPerStream;
    cudaMemcpyAsync(&d_a[offset], &h_a[offset], bytesPerStream, cudaMemcpyHostToDevice, streams[i]);
    kernel<<<grid, block, 0, streams[i]>>>(&d_a[offset]);
    cudaMemcpyAsync(&h_a[offset], &d_a[offset], bytesPerStream, cudaMemcpyDeviceToHost, streams[i]);
}

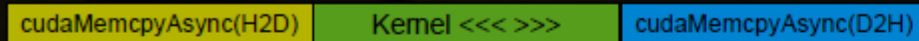
for (int i = 0; i < nStreams; i++) {
    cudaStreamsSynchronize(streams[i]);
}
```

Example: A Simple Timeline of CUDA Operations Using Three Streams



Number of Concurrency

- **Serial (1x)**



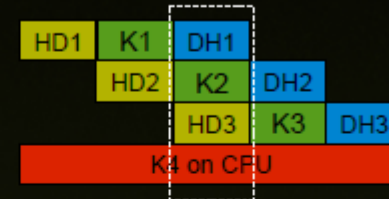
- **2-way concurrency (up to 2x)**



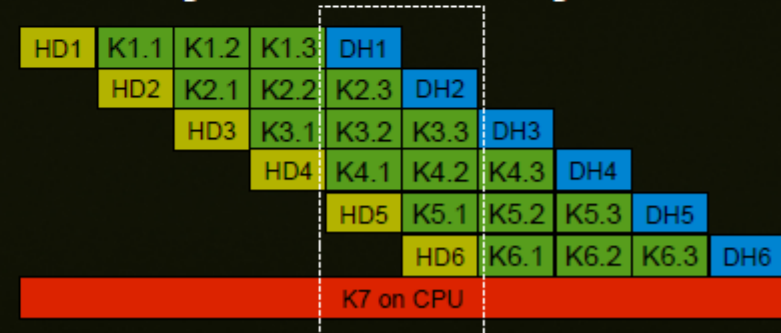
- **3-way concurrency (up to 3x)**



- **4-way concurrency (3x+)**



- **4+ way concurrency**



Conditions to Be Satisfied When Using Streams to Overlap Device Execution with Data Transfer

- First, the device must support a feature known as **device overlap**
 - A GPU supporting device overlap possesses the capacity to simultaneously execute a CUDA C kernel while performing a copy between device and host memory

```
cudaDeviceProp prop;
int whichDevice;
HANDLE_ERROR( cudaGetDevice( &whichDevice ) );
HANDLE_ERROR( cudaGetDeviceProperties( &prop, whichDevice ) );
if (!prop.deviceOverlap) {
    printf( "Device will not handle overlaps, so no speed up from streams\n" );
    return 0;
}
```

Conditions to Be Satisfied When Using Streams to Overlap Device Execution with Data Transfer

- Second, The kernel execution and the data transfer to be overlapped must both occur in different, non- NULL streams



Conditions to Be Satisfied When Using Streams to Overlap Device Execution with Data Transfer

- Third, the host memory involved in data transfer must be **pinned (page-locked, non-pageable) memory**:

- Cannot be swapped (paged) out by the OS
- Transferred using the host CPU
- Transferred using the direct memory access , can reach higher bandwidths for large transfers
- Has higher overhead for allocation

- Allocating pinned memory

- `cudaMallocHost()`
- `cudaHostAlloc()`

- Free allocated memory

- `cudaFreeHost()`

THREE TYPES OF MEMORY

▪ Device Memory

- Allocated using `cudaMalloc`
- Cannot be paged

▪ Pageable Host Memory

- Default allocation (e.g. `malloc`, `calloc`, `new`, etc)
- Can be paged in and out by the OS

▪ Pinned (Page-Locked) Host Memory

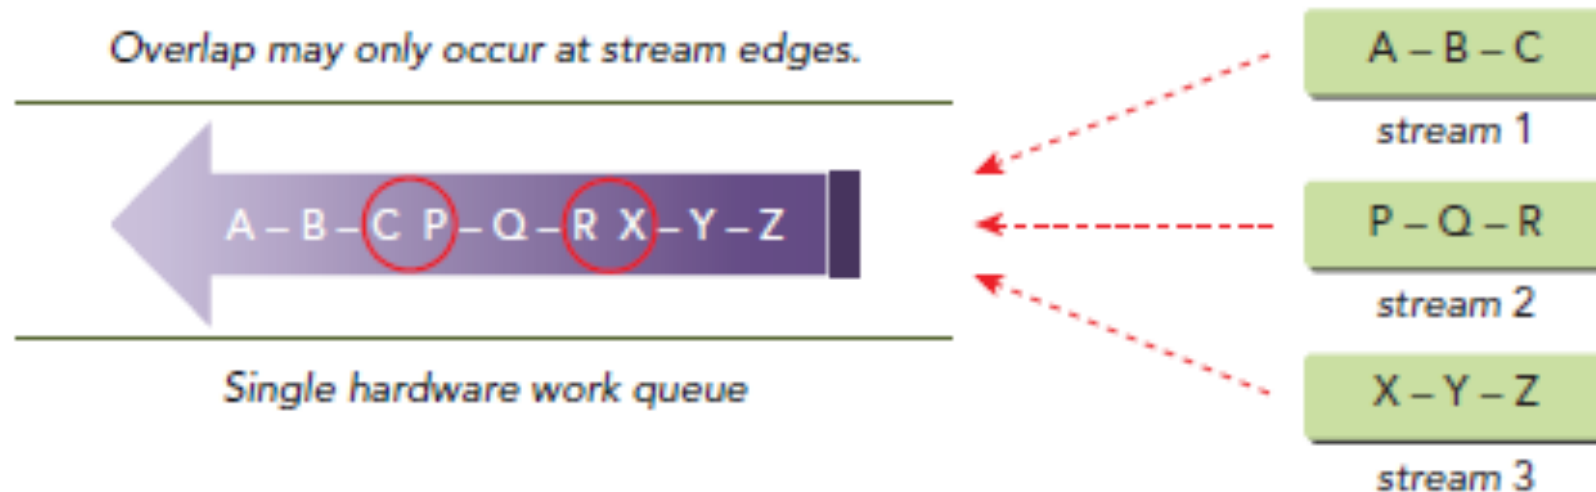
- Allocated using special allocators
- Cannot be paged out by the OS

Stream Scheduling

- Kernel and copy engine (possibly x2) have different queues
 - Fermi hardware has 3 queues
 - 1 Compute Engine queue
 - 2 Copy Engine queues – one for H2D and one for D2H
- CUDA operations are dispatched to hardware in sequence they were issued
 - CUDA operations are placed in the relevant queue
 - Stream dependencies between engine queues are maintained
 - Stream dependencies within an engine queue are lost
- A CUDA operation is dispatched from the engine queue if:
 - Preceding calls in the same stream have completed
 - Preceding calls in the same queue have been dispatched, and
 - Resources are available
- CUDA kernels may be executed concurrently if they are in different streams

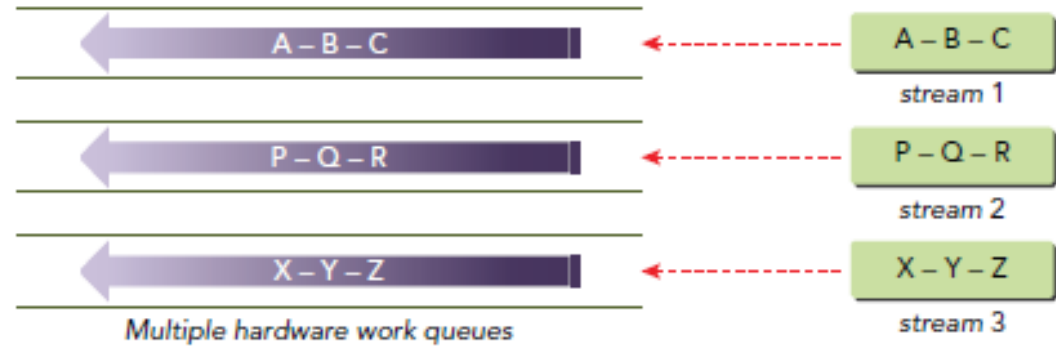
False Dependencies

- All streams are multiplexed into a single hardware work queue. The single pipeline may result in a false dependency for the preceding streams to block successive streams
- **Note a blocked operation blocks all other operations in the queue, even in other streams**



Hyper-Q

- False dependencies are reduced in the Kepler family of GPUs using multiple hardware work queues, a technology called **Hyper-Q**
- Hyper-Q allows multiple CPU threads or processes to launch work on a single GPU simultaneously by maintaining multiple hardware-managed connections between the host and the device
- `deviceProp.concurrentKernels`

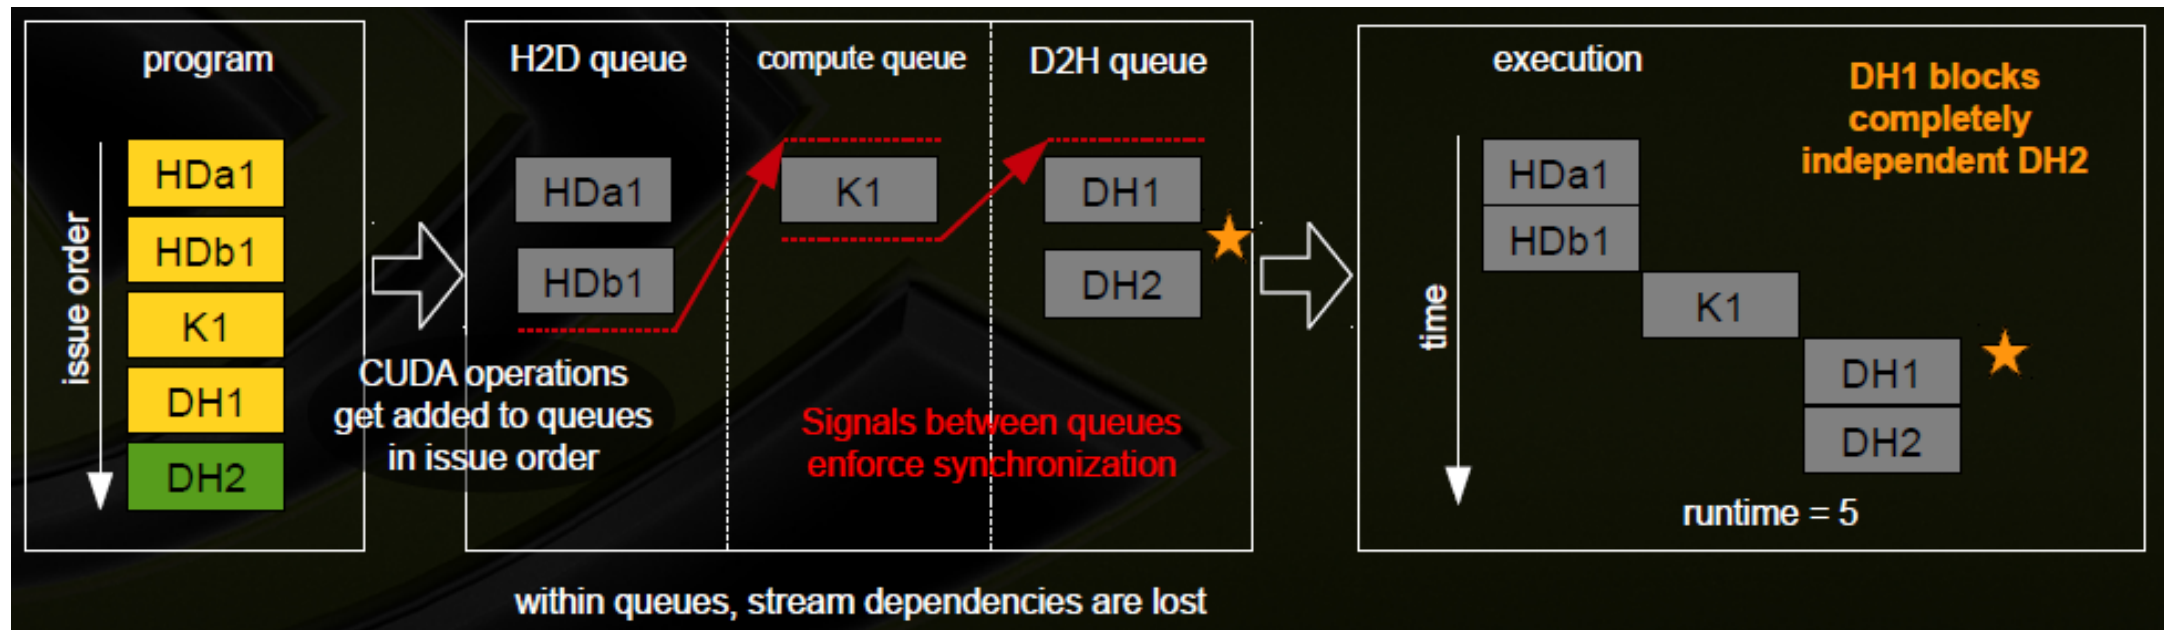


```
int dev = 0;
cudaDeviceProp deviceProp;
CHECK(cudaGetDeviceProperties(&deviceProp, dev));
printf("> Using Device %d: %s\n", dev, deviceProp.name);
CHECK(cudaSetDevice(dev));

// check if device support hyper-q
if (deviceProp.major < 3 || (deviceProp.major == 3 && deviceProp.minor < 5))
{
    if (deviceProp.concurrentKernels == 0)
    {
        printf("> GPU does not support concurrent kernel execution (SM 3.5 "
            "or higher required)\n");
        printf("> CUDA kernel runs will be serialized\n");
    }
    else
    {
        printf("> GPU does not support HyperQ\n");
        printf("> CUDA kernel runs will have limited concurrency\n");
    }
}
```

Example – Blocked Queue

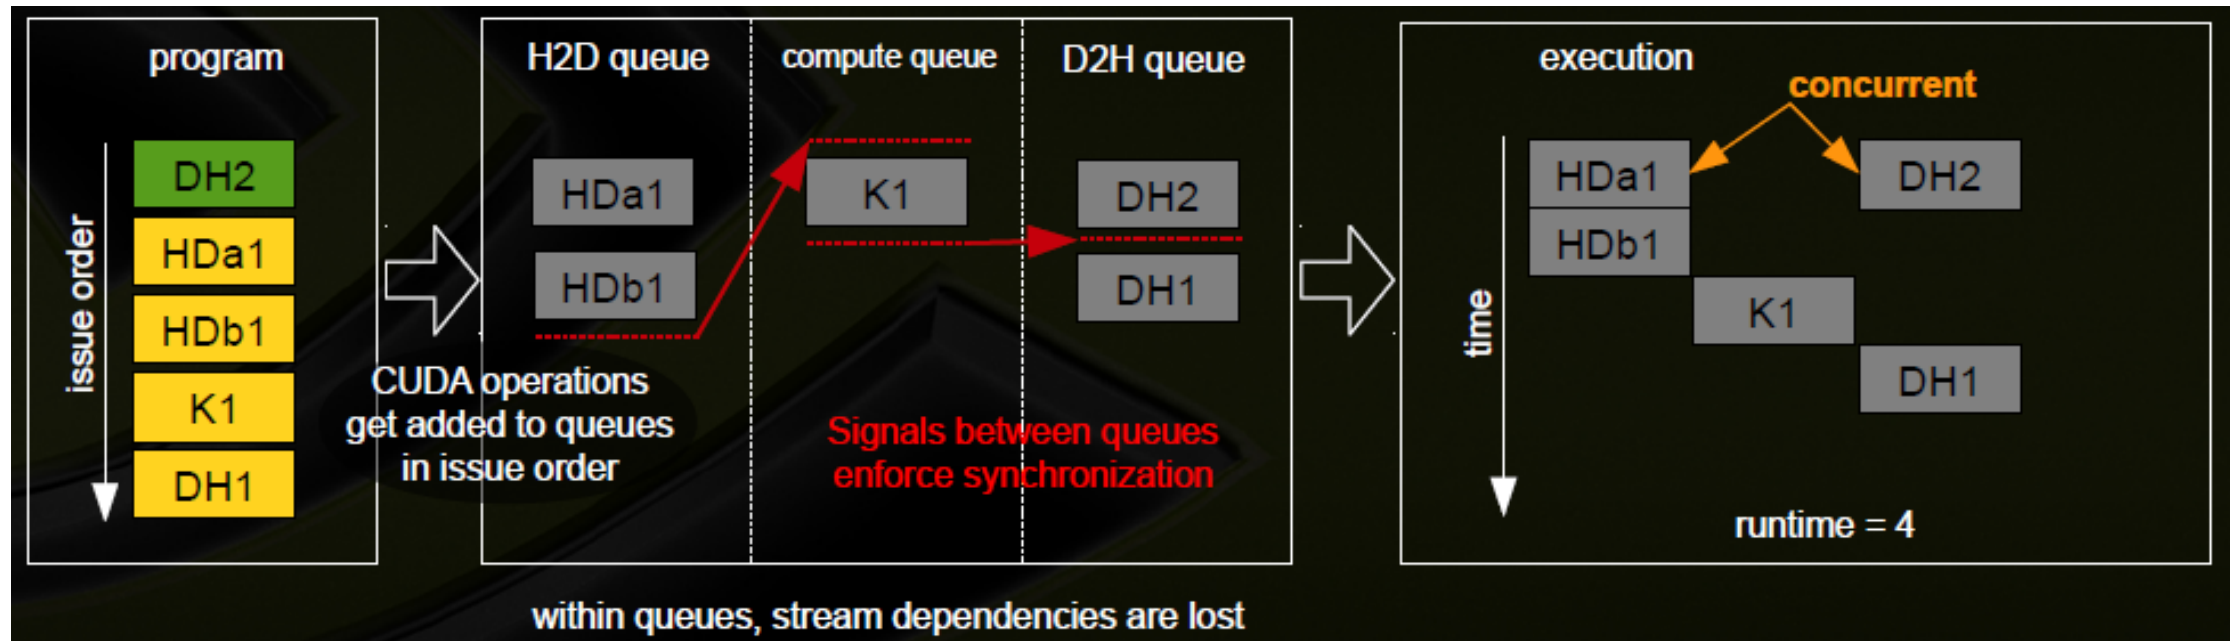
- Two streams, stream 1 is issued first
 - Stream 1: HDa1, HDb1, K1, DH1
 - Stream 2: DH2 (completely independent of stream 1)



Example – Blocked Queue

- Two streams, **stream 2 is issued first**
 - Stream 1: HDa1, HDb1, K1, DH1
 - Stream 2: DH2

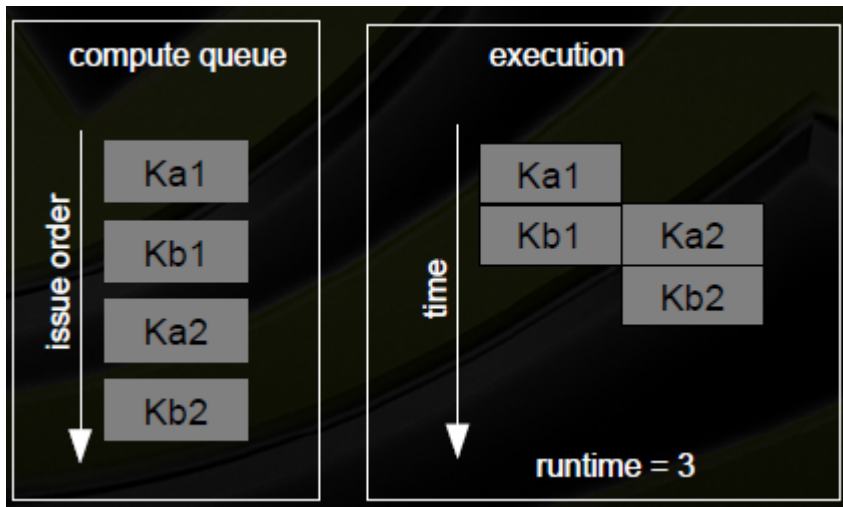
issue order matters!



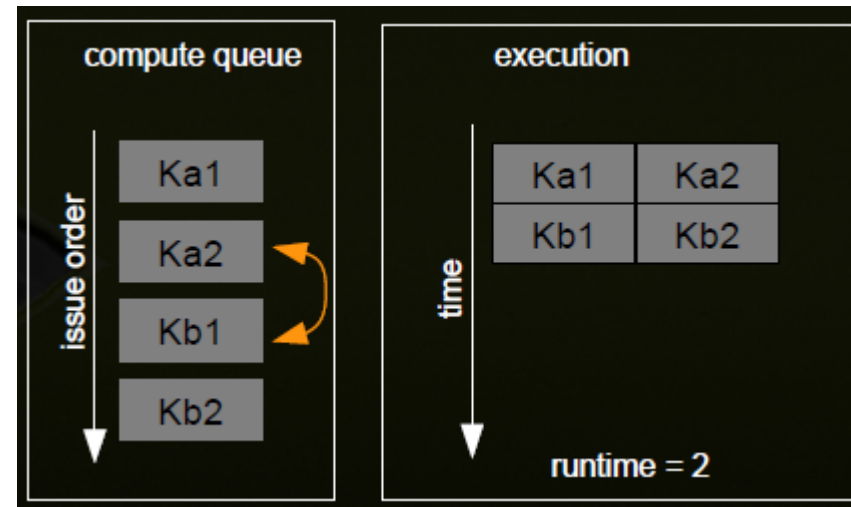
Example – Blocked Kernel

- Two streams – just issuing CUDA kernels
 - Stream 1: Ka1, Kb1
 - Stream 2: Ka2, Kb2
 - Kernels are **similar size**, fill $\frac{1}{2}$ of the SM resources

Issue depth first



Issue breadth first

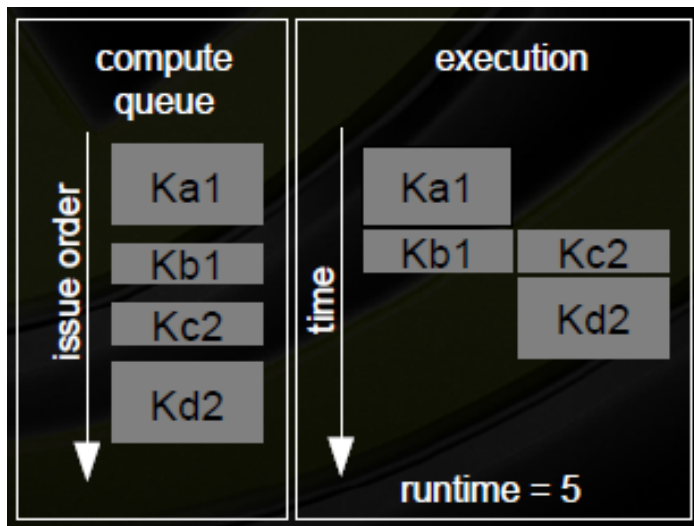


issue order matters!

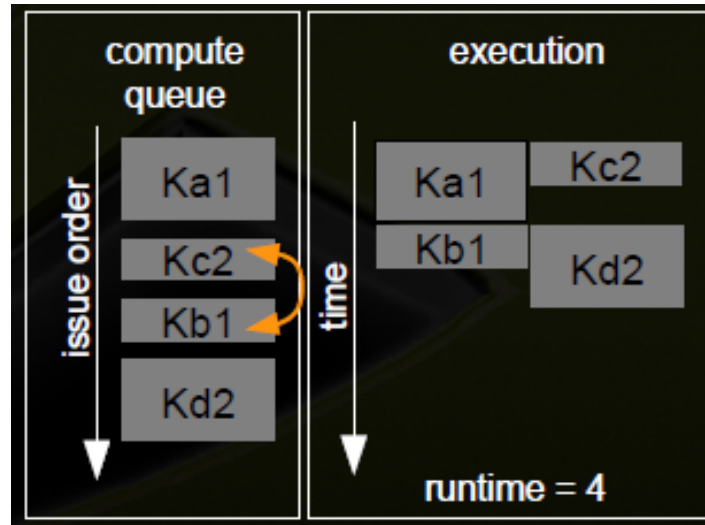
Example – Optimal Concurrency Can Depend on Kernel Execution Time

- Two streams – just issuing CUDA kernels – but kernels are **different “sizes”**
 - Stream 1: Ka1 {2}, Kb1 {1}
 - Stream 2: Kc2 {1}, Kd2 {2}
 - Kernels fill ½ of the SM resources

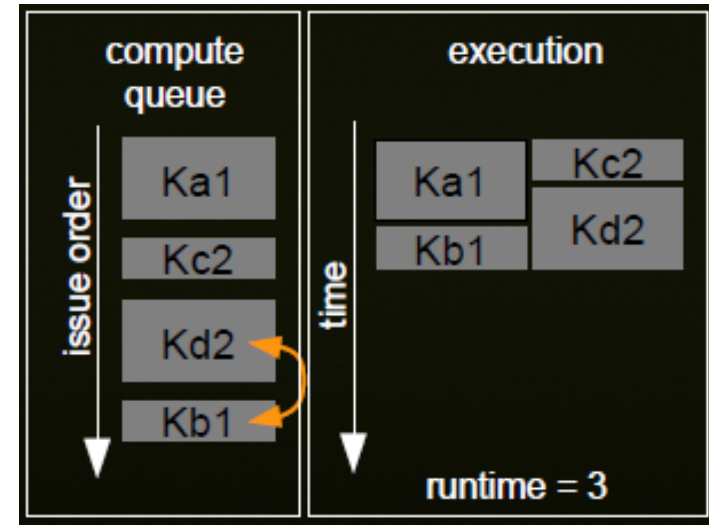
Issue depth first



Issue breadth first



Custom



**issue order matters!
execution time matters!**

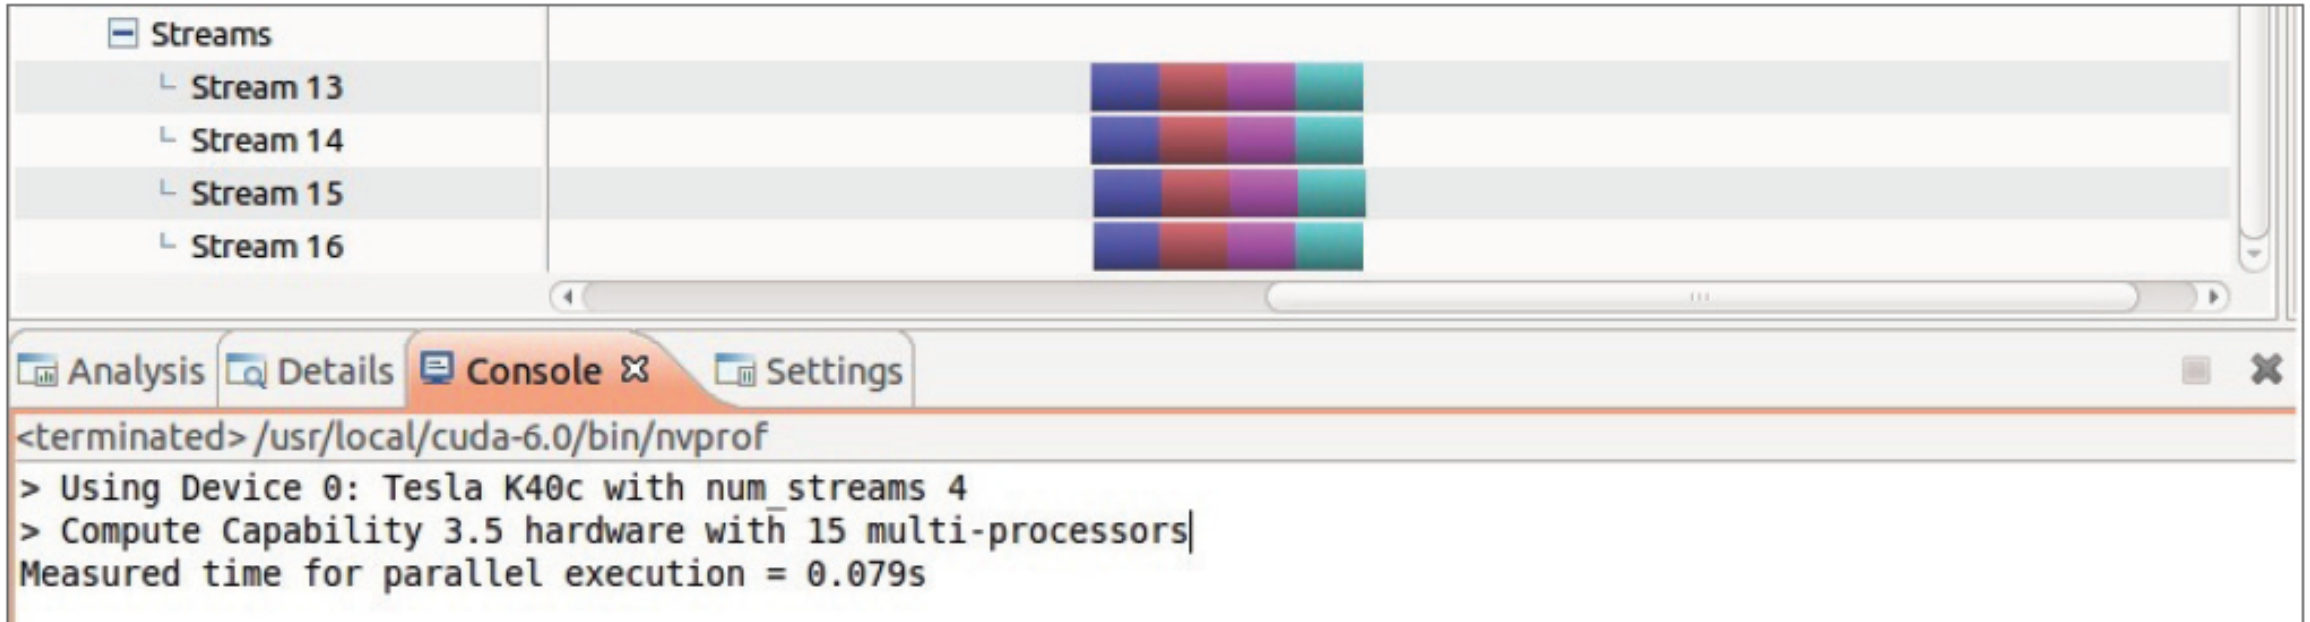
NVIDIA Visual Profiler (nvvp)

- Allows you to visualize and optimize the performance of your application
- Displays a timeline of your application's activity on both the CPU and GPU so that you can identify opportunities for performance improvement
- Analyzes your application to detect potential performance bottlenecks and direct you on how to take action to eliminate or reduce those bottlenecks
- **\$nvvp ./a.out**
- A timeline will contain a **Stream** row for each stream used by the application (including both the default stream and any application created streams). Each interval in a **Stream** row represents the duration of a memcpy or kernel execution performed on that stream

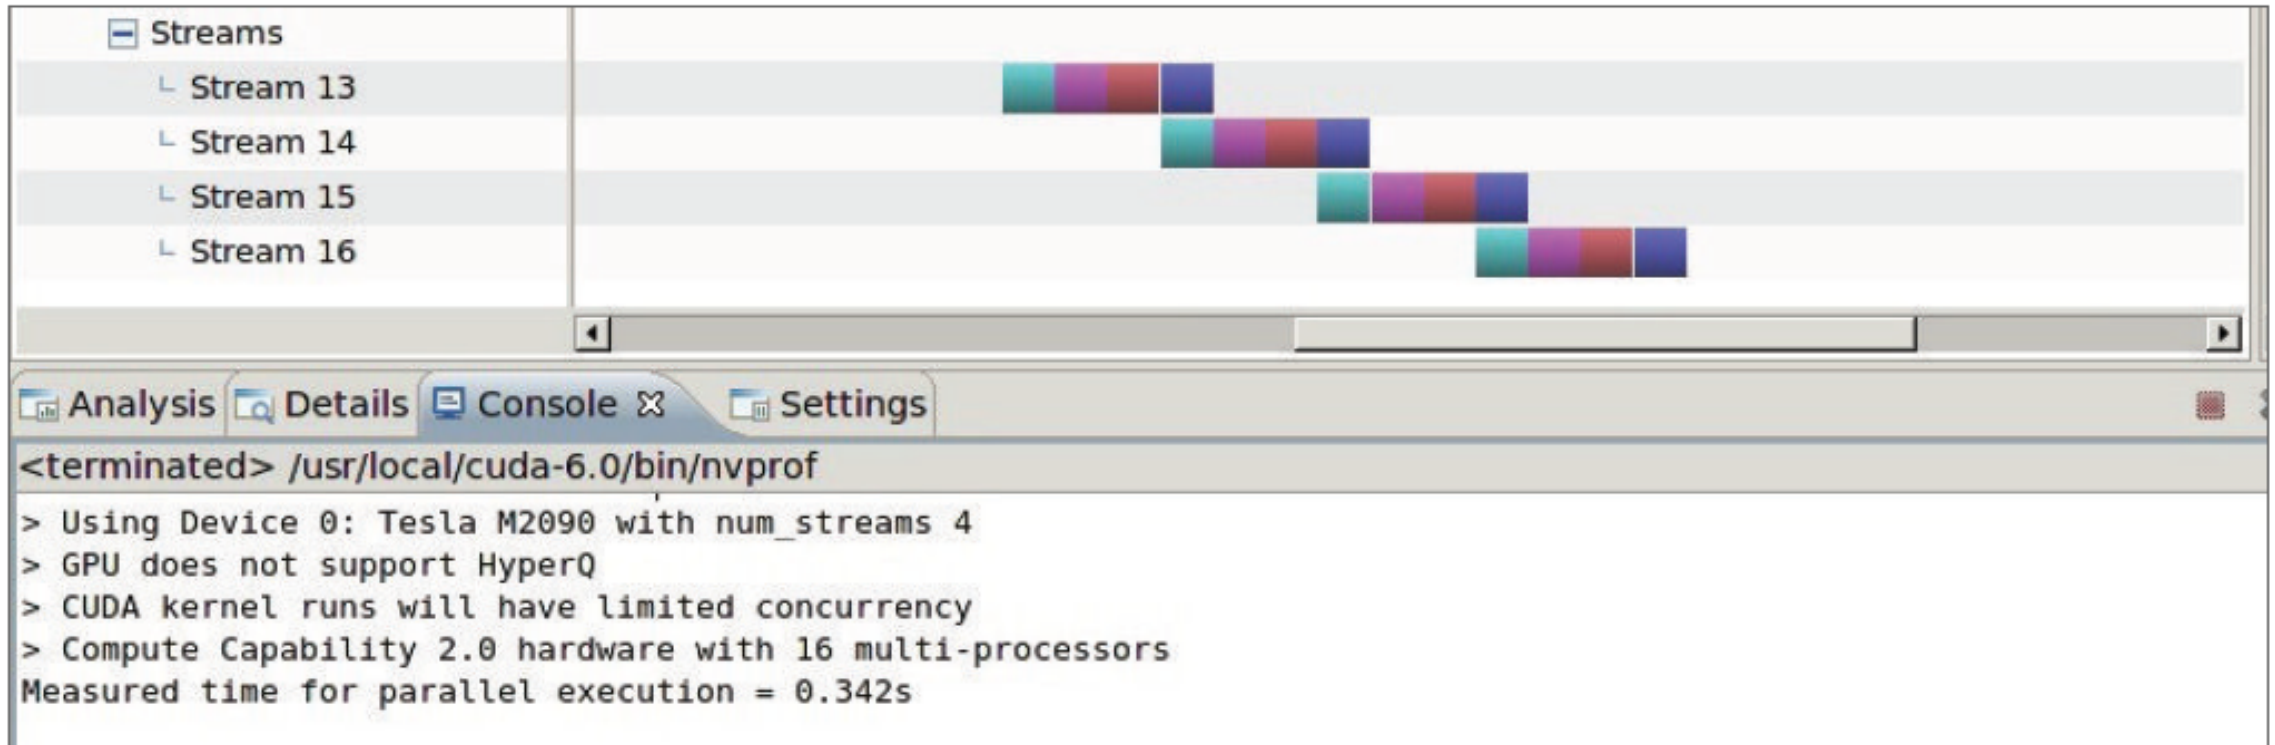
Visualizing Concurrent Kernel Executions on Tesla K40

```
for (int i = 0; i < n_streams; i++) {  
    kernel_1<<<grid, block, 0, streams[i]>>>();  
    kernel_2<<<grid, block, 0, streams[i]>>>();  
    kernel_3<<<grid, block, 0, streams[i]>>>();  
    kernel_4<<<grid, block, 0, streams[i]>>>();  
}
```

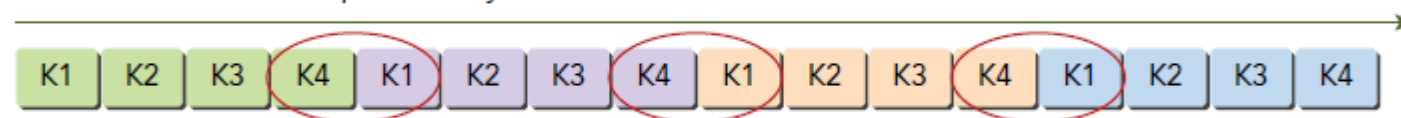
```
__global__ void kernel_1() {  
    double sum = 0.0;  
    for (int i = 0; i < N; i++) {  
        sum = sum + tan(0.1) * tan(0.1);  
    }  
}
```



Demonstrating False Dependencies with Depth-First Assignment on Fermi GPUs



Issue order from host: depth-first way



Only the three stream edges are independent.

Avoid False Dependencies on Fermi GPUs with Breadth-First Assignment

```
// dispatch job with breadth first way
for (int i = 0; i < n_streams; i++)
    kernel_1<<<grid, block, 0, streams[i]>>>();
for (int i = 0; i < n_streams; i++)
    kernel_2<<<grid, block, 0, streams[i]>>>();
for (int i = 0; i < n_streams; i++)
    kernel_3<<<grid, block, 0, streams[i]>>>();
for (int i = 0; i < n_streams; i++)
    kernel_4<<<grid, block, 0, streams[i]>>>();
```

Issue order from host: breadth-first order



There is no dependence between any adjacent kernels.

The screenshot shows the NVIDIA Visual Profiler interface. The left pane displays a list of streams: Stream 13, Stream 14, Stream 15, and Stream 16. The right pane shows a timeline view with four colored bars (green, purple, orange, blue) representing the execution of kernels K1, K2, K3, and K4 across the streams. The bottom pane shows the console output of the nvprof command.

Analysis Details Console Settings

```
<terminated> /usr/local/cuda-6.0/bin/nvprof
> Using Device 0: Tesla M2090 with num_streams 4
> GPU does not support HyperQ
> CUDA kernel runs will have limited concurrency
> Compute Capability 2.0 hardware with 16 multi-processors
Measured time for parallel execution = 0.105s
```

Coding Examples

- Pinned-memory
- Vector addition
 - Single stream version
 - Double stream with depth-first assignment version
 - Double stream with breadth-first assignment version

Example: Pinned Memory

- `cuda_malloc_test()` does the memory copy for pageable memory
 - Allocate host memory by `malloc()`
- `cuda_host_alloc_test()` does the memory copy for pinned memory
 - Allocate host memory by `cudaHostAlloc()`

```
float cuda_malloc_test( int size, bool up ) {
    cudaEvent_t    start, stop;
    int            *a, *dev_a;
    float          elapsedTime;

    HANDLE_ERROR( cudaEventCreate( &start ) );
    HANDLE_ERROR( cudaEventCreate( &stop ) );

    a = (int*)malloc( size * sizeof( *a ) );
    HANDLE_NULL( a );
    HANDLE_ERROR( cudaMalloc( (void**)&dev_a,
                             size * sizeof( *dev_a ) ) );
```

```
float cuda_host_alloc_test( int size, bool up ) {
    cudaEvent_t    start, stop;
    int            *a, *dev_a;
    float          elapsedTime;

    HANDLE_ERROR( cudaEventCreate( &start ) );
    HANDLE_ERROR( cudaEventCreate( &stop ) );

    HANDLE_ERROR( cudaHostAlloc( (void**)&a,
                                size * sizeof( *a ),
                                cudaHostAllocDefault ) );
    HANDLE_ERROR( cudaMalloc( (void**)&dev_a,
                             size * sizeof( *dev_a ) ) );
```

Example: Pinned Memory

- The loops of copy, and stop events for timer are the same inside both `cuda_malloc_test()` and `cuda_host_alloc_test()`

```
HANDLE_ERROR( cudaEventRecord( start, 0 ) );
for (int i=0; i<100; i++) {
    if (up)
        HANDLE_ERROR( cudaMemcpy( dev_a, a,
                                   size * sizeof( *dev_a ),
                                   cudaMemcpyHostToDevice ) );
    else
        HANDLE_ERROR( cudaMemcpy( a, dev_a,
                                   size * sizeof( *dev_a ),
                                   cudaMemcpyDeviceToHost ) );
}
HANDLE_ERROR( cudaEventRecord( stop, 0 ) );
HANDLE_ERROR( cudaEventSynchronize( stop ) );
HANDLE_ERROR( cudaEventElapsedTime( &elapsedTime,
                                   start, stop ) );
```

Example: Pinned Memory

- `cuda_malloc_test()` frees memory by `free()`
- `cuda_host_alloc_test()` frees memory by `cudaFreeHost()`

```
free( a );  
HANDLE_ERROR( cudaFree( dev_a ) );  
HANDLE_ERROR( cudaEventDestroy( start ) );  
HANDLE_ERROR( cudaEventDestroy( stop ) );  
  
return elapsedTime;  
}
```

```
HANDLE_ERROR( cudaFreeHost( a ) );  
HANDLE_ERROR( cudaFree( dev_a ) );  
HANDLE_ERROR( cudaEventDestroy( start ) );  
HANDLE_ERROR( cudaEventDestroy( stop ) );  
  
return elapsedTime;  
}
```


Example: Pinned Memory

```
#include "../common/book.h"

#define SIZE      (64*1024*1024)

...

int main( void ) {
    float elapsedTime;
    float MB = (float)100*SIZE*sizeof(int)/1024/1024;

    // try it with cudaMalloc
    elapsedTime = cuda_malloc_test( SIZE, true );
    printf( "Time using cudaMalloc:   %3.1f ms\n", elapsedTime );
    printf( "\tMB/s during copy up:   %3.1f\n", MB/(elapsedTime/1000) );

    elapsedTime = cuda_malloc_test( SIZE, false );
    printf( "Time using cudaMalloc:   %3.1f ms\n", elapsedTime );
    printf( "\tMB/s during copy down:   %3.1f\n", MB/(elapsedTime/1000) );

    // now try it with cudaHostAlloc
    elapsedTime = cuda_host_alloc_test( SIZE, true );
    printf( "Time using cudaHostAlloc: %3.1f ms\n", elapsedTime );
    printf( "\tMB/s during copy up:   %3.1f\n", MB/(elapsedTime/1000) );

    elapsedTime = cuda_host_alloc_test( SIZE, false );
    printf( "Time using cudaHostAlloc: %3.1f ms\n", elapsedTime );
    printf( "\tMB/s during copy down: %3.1f\n", MB/(elapsedTime/1000) );
}
```

```
[[jin6@node1733 11_StreamsBasics]$ ./a.out
Time using cudaMalloc:  6688.8 ms
                        MB/s during copy up:  3827.3
Time using cudaMalloc:  7377.6 ms
                        MB/s during copy down:  3470.0
Time using cudaHostAlloc:  4281.8 ms
                        MB/s during copy up:  5978.7
Time using cudaHostAlloc:  4050.1 ms
                        MB/s during copy down:  6320.8
```


Example: Vector Addition

- Use a CUDA kernel to take two input buffers of data, a and b
- The kernel compute an average of three values in a and b to produce an output buffer c

```
17  #include "../common/book.h"
18
19  #define N      (1024*1024)
20  #define FULL_DATA_SIZE      (N*20)
21
22
23  __global__ void kernel( int *a, int *b, int *c ) {
24      int idx = threadIdx.x + blockIdx.x * blockDim.x;
25      if (idx < N) {
26          int idx1 = (idx + 1) % 256;
27          int idx2 = (idx + 2) % 256;
28          float  as = (a[idx] + a[idx1] + a[idx2]) / 3.0f;
29          float  bs = (b[idx] + b[idx1] + b[idx2]) / 3.0f;
30          c[idx] = (as + bs) / 2;
31      }
32  }
```

Example: Vector Addition – Single Stream

- Device check and declarations

```
35  int main( void ) {
36      cudaDeviceProp  prop;
37      int whichDevice;
38      HANDLE_ERROR( cudaGetDevice( &whichDevice ) );
39      HANDLE_ERROR( cudaGetDeviceProperties( &prop, whichDevice ) );
40      if ( !prop.deviceOverlap ) {
41          printf( "Device will not handle overlaps, so no speed up from streams\n" );
42          return 0;
43      }
44
45      cudaEvent_t      start, stop;
46      float            elapsedTime;
47
48      cudaStream_t      stream;
49      int *host_a, *host_b, *host_c;
50      int *dev_a, *dev_b, *dev_c;
51
52      // start the timers
53      HANDLE_ERROR( cudaEventCreate( &start ) );
54      HANDLE_ERROR( cudaEventCreate( &stop ) );
55
56      // initialize the stream
57      HANDLE_ERROR( cudaStreamCreate( &stream ) );
```

Example: Vector Addition – Single Stream

- Memory allocations and initializations

```
59 // allocate the memory on the GPU
60 HANDLE_ERROR( cudaMalloc( (void**)&dev_a,
61                           N * sizeof(int) ) );
62 HANDLE_ERROR( cudaMalloc( (void**)&dev_b,
63                           N * sizeof(int) ) );
64 HANDLE_ERROR( cudaMalloc( (void**)&dev_c,
65                           N * sizeof(int) ) );
66
67 // allocate host locked memory, used to stream
68 HANDLE_ERROR( cudaHostAlloc( (void**)&host_a,
69                              FULL_DATA_SIZE * sizeof(int),
70                              cudaHostAllocDefault ) );
71 HANDLE_ERROR( cudaHostAlloc( (void**)&host_b,
72                              FULL_DATA_SIZE * sizeof(int),
73                              cudaHostAllocDefault ) );
74 HANDLE_ERROR( cudaHostAlloc( (void**)&host_c,
75                              FULL_DATA_SIZE * sizeof(int),
76                              cudaHostAllocDefault ) );
77
78 for (int i=0; i<FULL_DATA_SIZE; i++) {
79     host_a[i] = rand();
80     host_b[i] = rand();
81 }
```

Example: Vector Addition – Single Stream

- Split data, perform kernel operations, and copy result back

```
83     HANDLE_ERROR( cudaEventRecord( start, 0 ) );
84     // now loop over full data, in bite-sized chunks
85     for (int i=0; i<FULL_DATA_SIZE; i+= N) {
86         // copy the locked memory to the device, async
87         HANDLE_ERROR( cudaMemcpyAsync( dev_a, host_a+i,
88                                     N * sizeof(int),
89                                     cudaMemcpyHostToDevice,
90                                     stream ) );
91         HANDLE_ERROR( cudaMemcpyAsync( dev_b, host_b+i,
92                                     N * sizeof(int),
93                                     cudaMemcpyHostToDevice,
94                                     stream ) );
95
96         kernel<<<N/256,256,0,stream>>>( dev_a, dev_b, dev_c );
97
98         // copy the data from device to locked memory
99         HANDLE_ERROR( cudaMemcpyAsync( host_c+i, dev_c,
100                                     N * sizeof(int),
101                                     cudaMemcpyDeviceToHost,
102                                     stream ) );
103
104     }
105     // copy result chunk from locked to full buffer
106     HANDLE_ERROR( cudaStreamSynchronize( stream ) );
```

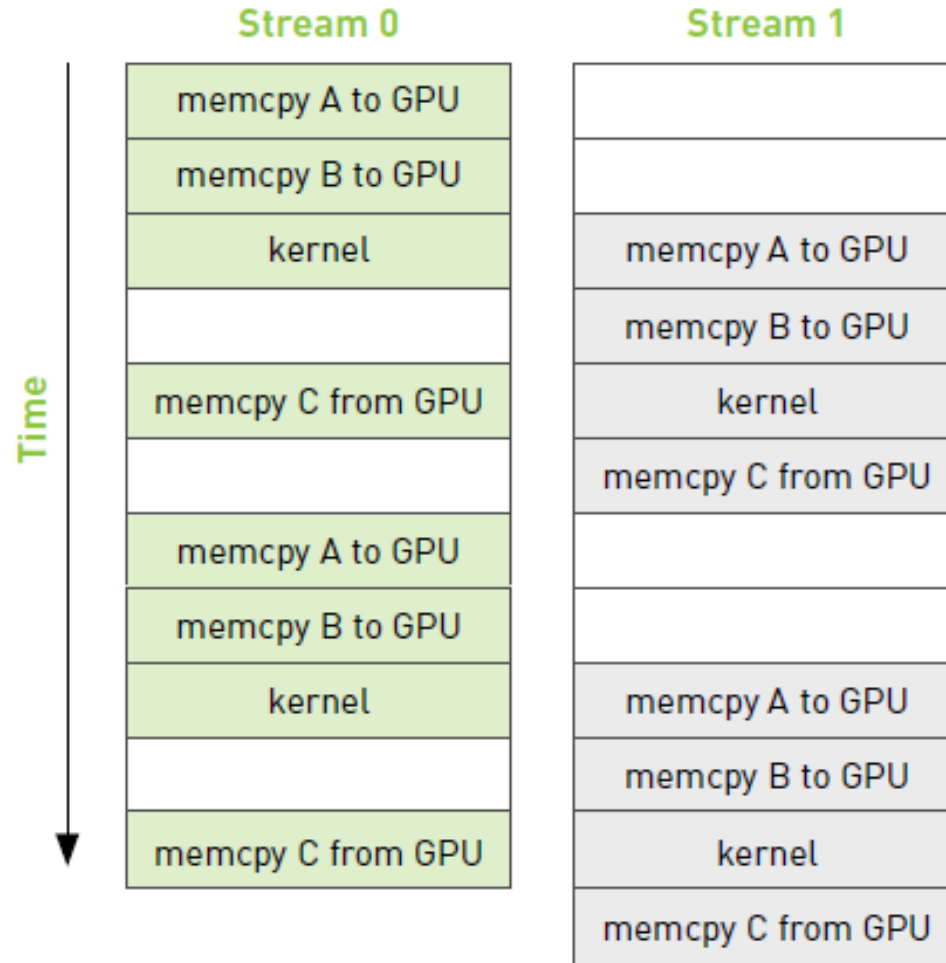
Example: Vector Addition – Single Stream

- Stop timer, collect performance data, free buffers, and destroy stream

```
108     HANDLE_ERROR( cudaEventRecord( stop, 0 ) );
109
110     HANDLE_ERROR( cudaEventSynchronize( stop ) );
111     HANDLE_ERROR( cudaEventElapsedTime( &elapsedTime,
112                                         start, stop ) );
113     printf( "Time taken:  %3.1f ms\n", elapsedTime );
114
115     // cleanup the streams and memory
116     HANDLE_ERROR( cudaFreeHost( host_a ) );
117     HANDLE_ERROR( cudaFreeHost( host_b ) );
118     HANDLE_ERROR( cudaFreeHost( host_c ) );
119     HANDLE_ERROR( cudaFree( dev_a ) );
120     HANDLE_ERROR( cudaFree( dev_b ) );
121     HANDLE_ERROR( cudaFree( dev_c ) );
122     HANDLE_ERROR( cudaStreamDestroy( stream ) );
123
124     return 0;
125 }
```

Example: Vector Addition – Double Streams

- The idea underlying this version relies on two things:
 - The “chunked” computation, and
 - The overlap of memory copies with kernel execution.
- Enqueue operations across streams
 - Depth-first
 - Breadth-first



Timeline of intended application execution using two independent streams. (Calls to `cudaMemcpyAsync()` are abbreviated to “memcpy”.)

Example: Vector Addition – Double Streams (Depth-first)

```
35 int main( void ) {
36     cudaDeviceProp prop;
37     int whichDevice;
38     HANDLE_ERROR( cudaGetDevice( &whichDevice ) );
39     HANDLE_ERROR( cudaGetDeviceProperties( &prop, whichDevice ) );
40     if (!prop.deviceOverlap) {
41         printf( "Device will not handle overlaps, so no speed up from streams\n" );
42         return 0;
43     }
44
45     cudaEvent_t      start, stop;
46     float            elapsedTime;
47
48     cudaStream_t      stream0, stream1;
49     int *host a, *host b, *host c;
50     int *dev_a0, *dev_b0, *dev_c0;
51     int *dev_a1, *dev_b1, *dev_c1;
52
53     // start the timers
54     HANDLE_ERROR( cudaEventCreate( &start ) );
55     HANDLE_ERROR( cudaEventCreate( &stop ) );
56
57     // initialize the streams
58     HANDLE_ERROR( cudaStreamCreate( &stream0 ) );
59     HANDLE_ERROR( cudaStreamCreate( &stream1 ) );
```

Example: Vector Addition – Double Streams (Depth-first)

```
61 // allocate the memory on the GPU
62 HANDLE_ERROR( cudaMalloc( (void**)&dev_a0,
63                           N * sizeof(int) ) );
64 HANDLE_ERROR( cudaMalloc( (void**)&dev_b0,
65                           N * sizeof(int) ) );
66 HANDLE_ERROR( cudaMalloc( (void**)&dev_c0,
67                           N * sizeof(int) ) );
68 HANDLE_ERROR( cudaMalloc( (void**)&dev_a1,
69                           N * sizeof(int) ) );
70 HANDLE_ERROR( cudaMalloc( (void**)&dev_b1,
71                           N * sizeof(int) ) );
72 HANDLE_ERROR( cudaMalloc( (void**)&dev_c1,
73                           N * sizeof(int) ) );
74
75 // allocate host locked memory, used to stream
76 HANDLE_ERROR( cudaHostAlloc( (void**)&host_a,
77                              FULL_DATA_SIZE * sizeof(int),
78                              cudaHostAllocDefault ) );
79 HANDLE_ERROR( cudaHostAlloc( (void**)&host_b,
80                              FULL_DATA_SIZE * sizeof(int),
81                              cudaHostAllocDefault ) );
82 HANDLE_ERROR( cudaHostAlloc( (void**)&host_c,
83                              FULL_DATA_SIZE * sizeof(int),
84                              cudaHostAllocDefault ) );
85
86 for (int i=0; i<FULL_DATA_SIZE; i++) {
87     host_a[i] = rand();
88     host_b[i] = rand();
89 }
```

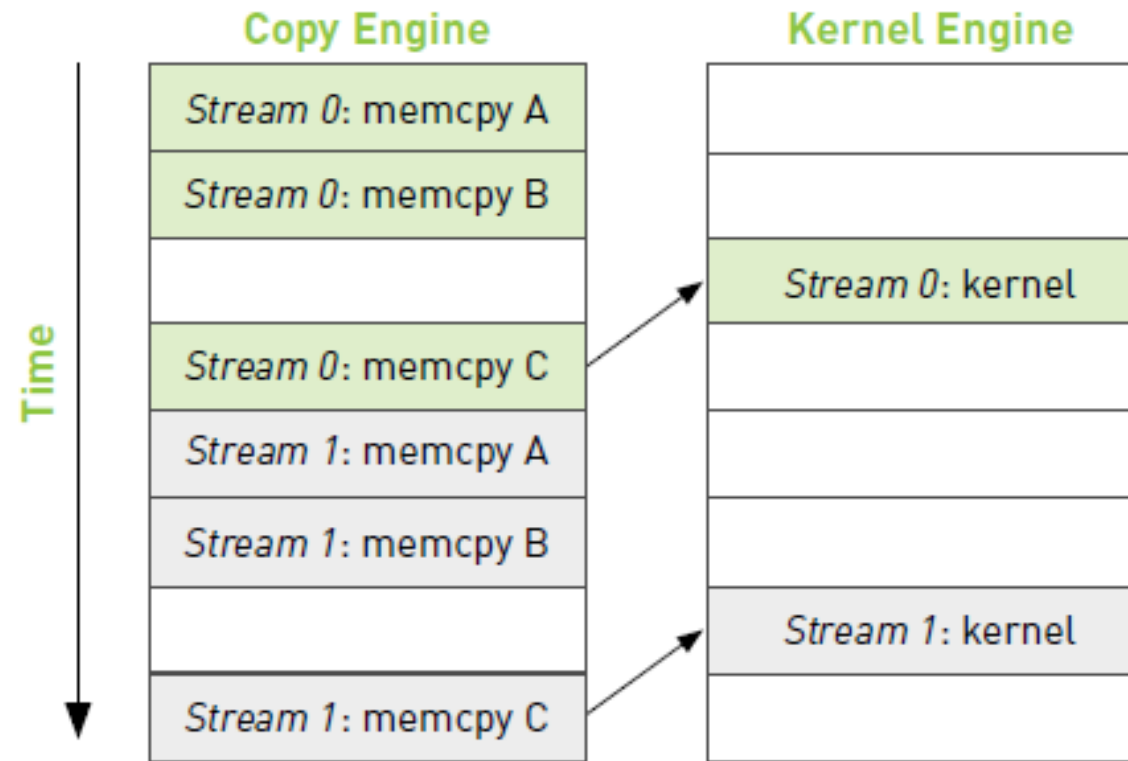

Example: Vector Addition – Double Streams (Depth-first)

```
91 HANDLE_ERROR( cudaEventRecord( start, 0 ) );
92 // now loop over full data, in bite-sized chunks
93 for (int i=0; i<FULL_DATA_SIZE; i+= N*2) {
94     // copy the locked memory to the device, async
95     HANDLE_ERROR( cudaMemcpyAsync( dev_a0, host_a+i, N * sizeof(int), cudaMemcpyHostToDevice, stream0 ) );
96     HANDLE_ERROR( cudaMemcpyAsync( dev_b0, host_b+i, N * sizeof(int), cudaMemcpyHostToDevice, stream0 ) );
97
98     kernel<<<N/256,256,0,stream0>>>( dev_a0, dev_b0, dev_c0 );
99
100    // copy the data from device to locked memory
101    HANDLE_ERROR( cudaMemcpyAsync( host_c+i, dev_c0, N * sizeof(int), cudaMemcpyDeviceToHost, stream0 ) );
102
103    // copy the locked memory to the device, async
104    HANDLE_ERROR( cudaMemcpyAsync( dev_a1, host_a+i+N, N * sizeof(int), cudaMemcpyHostToDevice, stream1 ) );
105    HANDLE_ERROR( cudaMemcpyAsync( dev_b1, host_b+i+N, N * sizeof(int), cudaMemcpyHostToDevice, stream1 ) );
106
107    kernel<<<N/256,256,0,stream1>>>( dev_a1, dev_b1, dev_c1 );
108
109    // copy the data from device to locked memory
110    HANDLE_ERROR( cudaMemcpyAsync( host_c+i+N, dev_c1, N * sizeof(int), cudaMemcpyDeviceToHost, stream1 ) );
111 }
112 HANDLE_ERROR( cudaStreamSynchronize( stream0 ) );
113 HANDLE_ERROR( cudaStreamSynchronize( stream1 ) );
```

Example: Vector Addition – Double Streams (Depth-first)

```
115     HANDLE_ERROR( cudaEventRecord( stop, 0 ) );
116
117     HANDLE_ERROR( cudaEventSynchronize( stop ) );
118     HANDLE_ERROR( cudaEventElapsedTime( &elapsedTime,
119                                         start, stop ) );
120     printf( "Time taken:  %3.1f ms\n", elapsedTime );
121
122     // cleanup the streams and memory
123     HANDLE_ERROR( cudaFreeHost( host_a ) );
124     HANDLE_ERROR( cudaFreeHost( host_b ) );
125     HANDLE_ERROR( cudaFreeHost( host_c ) );
126     HANDLE_ERROR( cudaFree( dev_a0 ) );
127     HANDLE_ERROR( cudaFree( dev_b0 ) );
128     HANDLE_ERROR( cudaFree( dev_c0 ) );
129     HANDLE_ERROR( cudaFree( dev_a1 ) );
130     HANDLE_ERROR( cudaFree( dev_b1 ) );
131     HANDLE_ERROR( cudaFree( dev_c1 ) );
132     HANDLE_ERROR( cudaStreamDestroy( stream0 ) );
133     HANDLE_ERROR( cudaStreamDestroy( stream1 ) );
134
135     return 0;
136 }
```

Example: Vector Addition – Double Streams (Depth-first)

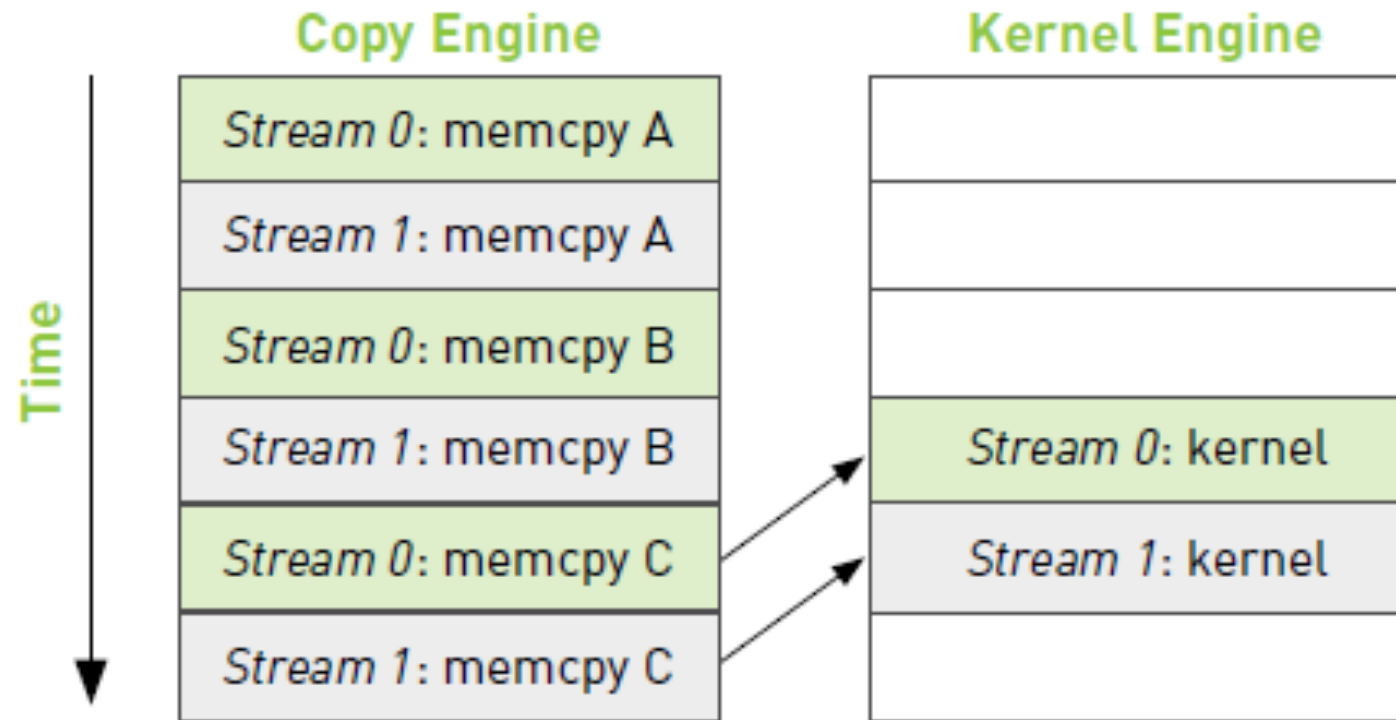


Stream 0's copy of c back to the host depends on its kernel execution completing
Stream 1's completely independent copies of a and b to the GPU get blocked because the GPU's engines execute work in the order it's provided
(Arrow's depicting the dependency of `cudaMemcpyAsync()` calls on kernel executions)

Example: Vector Addition – Double Streams (Breadth-first)

```
92 // now loop over full data, in bite-sized chunks
93 for (int i=0; i<FULL_DATA_SIZE; i+= N*2) {
94     // enqueue copies of a in stream0 and stream1
95     HANDLE_ERROR( cudaMemcpyAsync( dev_a0, host_a+i, N * sizeof(int), cudaMemcpyHostToDevice, stream0 ) );
96     HANDLE_ERROR( cudaMemcpyAsync( dev_a1, host_a+i+N, N * sizeof(int), cudaMemcpyHostToDevice, stream1 ) );
97     // enqueue copies of b in stream0 and stream1
98     HANDLE_ERROR( cudaMemcpyAsync( dev_b0, host_b+i, N * sizeof(int), cudaMemcpyHostToDevice, stream0 ) );
99     HANDLE_ERROR( cudaMemcpyAsync( dev_b1, host_b+i+N, N * sizeof(int), cudaMemcpyHostToDevice, stream1 ) );
100
101     // enqueue kernels in stream0 and stream1
102     kernel<<<N/256,256,0,stream0>>>( dev_a0, dev_b0, dev_c0 );
103     kernel<<<N/256,256,0,stream1>>>( dev_a1, dev_b1, dev_c1 );
104
105     // enqueue copies of c from device to locked memory
106     HANDLE_ERROR( cudaMemcpyAsync( host_c+i, dev_c0, N * sizeof(int), cudaMemcpyDeviceToHost, stream0 ) );
107     HANDLE_ERROR( cudaMemcpyAsync( host_c+i+N, dev_c1, N * sizeof(int), cudaMemcpyDeviceToHost, stream1 ) );
108 }
109 HANDLE_ERROR( cudaStreamSynchronize( stream0 ) );
110 HANDLE_ERROR( cudaStreamSynchronize( stream1 ) );
```

Example: Vector Addition – Double Streams (Breadth-first)



Example: Vector Addition – Results

```
[jin6@node1733 11_Streams]$ ./basic_single_stream  
Time taken: 48.8 ms  
[jin6@node1733 11_Streams]$ ./basic_double_stream_depth  
Time taken: 31.0 ms  
[jin6@node1733 11_Streams]$ ./basic_double_stream_breadth  
Time taken: 31.1 ms
```

Conclusions

- Using two or more CUDA streams enables simultaneous execution of kernel operations
- Asynchronous functions need to be performed with pinned memory allocated by `cudaHostAlloc()`
- The order in which we add operations to streams affect the capacity to achieve overlapping of memory copies and kernel executions
 - A general guideline is to use breadth-first assignment
- nvvp allows to you connect and view profiling data visually