

CPSC/ECE 4780/6780

General-Purpose Computation on Graphical Processing Units (GPGPU)

Lecture 11: Case Study: Advanced MRI Reconstruction

Recap of Last Lecture

- Why do we use multiple GPUs?
- What's the connectivity of multi-GPU system?
- What kind of inter-GPU communication patterns do we have?
- How to manage multiple GPUs from a single CPU thread?
- What is CUDA P2P APIs?
- What's the typical workflow for using streams and events in a multi-GPU application?

Case Study: Advanced MRI Reconstruction

1 an Application Case Study

- magnetic resonance imaging
- iterative reconstruction

2 Acceleration on GPU

- determining the kernel parallelism structure
- loop splitting
- loop interchange
- using registers to reduce memory accesses
- chunking data to fit into constant memory
- using hardware trigonometry functions

MCS 572 Lecture 38
Introduction to Supercomputing
Jan Verschelde, 18 November 2016

Advanced MRI Reconstruction

1 an Application Case Study

- magnetic resonance imaging
- iterative reconstruction

2 Acceleration on GPU

- determining the kernel parallelism structure
- loop splitting
- loop interchange
- using registers to reduce memory accesses
- chunking data to fit into constant memory
- using hardware trigonometry functions

magnetic resonance imaging

Magnetic Resonance Imaging (MRI) is a safe and noninvasive probe of the structure and function of tissues in the body.

MRI consists of two phases:

- 1 Acquisition or scan: the scanner samples data in the spatial-frequency domain along a predefined trajectory.
- 2 Reconstruction of the samples into an image.

Limitations: noise, imaging artifacts, long acquisition times.

Three often conflicting goals:

- Short scan time to reduce patient discomfort.
- High resolution and fidelity for early detection.
- High signal-to-noise ratio (SNR).

Massively parallel computing provides disruptive breakthrough.

problem formulation

The reconstructed image $m(\mathbf{r})$ is

$$\hat{m}(\mathbf{r}) = \sum_j W(\mathbf{k}_j) s(\mathbf{k}_j) e^{i2\pi \mathbf{k}_j \cdot \mathbf{r}}$$

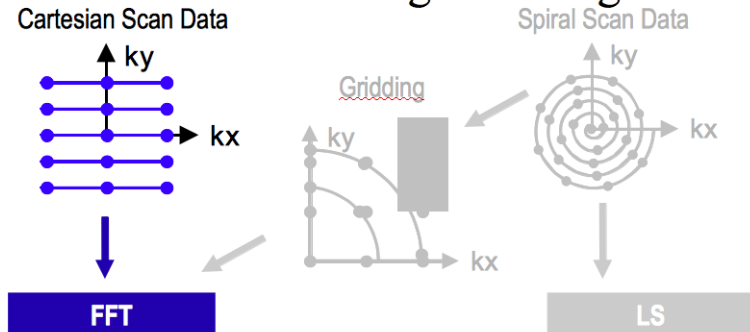
where

- $W(\mathbf{k})$ is the weighting function to account for nonuniform sampling;
- $s(\mathbf{k})$ is the measured k -space data.

The reconstruction is an inverse fast Fourier Transform on $s(\mathbf{k})$.

Cartesian trajectory with FFT reconstruction

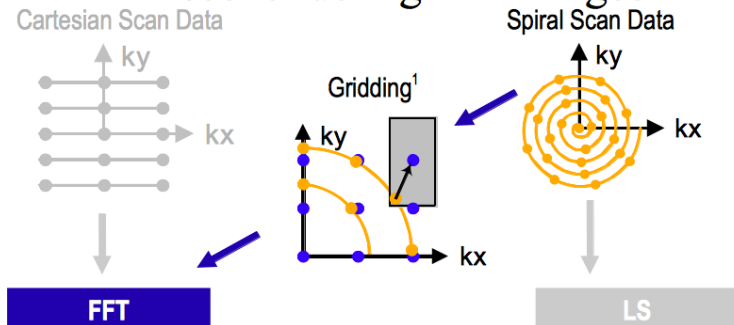
Reconstructing MR Images



Cartesian scan data + FFT:
Slow scan, fast reconstruction, images may be poor

spiral trajectory, gridding to enable FFT

Reconstructing MR Images



**Spiral scan data + Gridding + FFT:
Fast scan, fast reconstruction, better images**

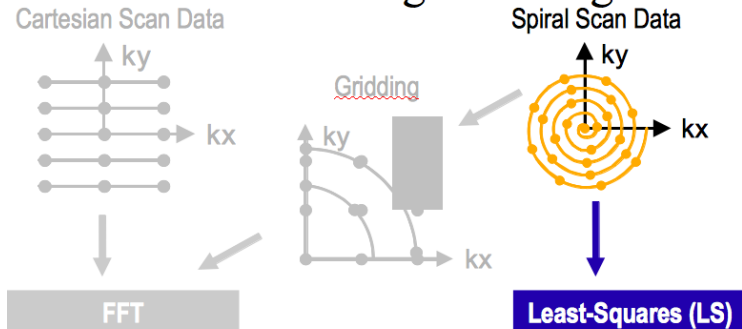
¹ Based on Fig 1 of Lustig et al, Fast Spiral Fourier Transform for Iterative MR Image Reconstruction, IEEE Int'l Symp. on Biomedical Imaging, 2004

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2009
University of Illinois, Urbana-Champaign

8

spiral trajectory with linear solver reconstruction

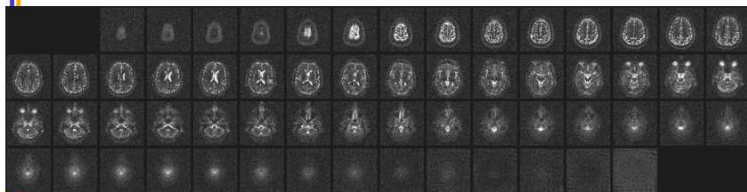
Reconstructing MR Images



Spiral scan data + LS
Superior images at expense of significantly more computation

sodium is much less abundant than water

An Exciting Revolution - Sodium Map of



- Images of sodium in the brain
 - Very large number of samples for increased SNR
 - Requires high-quality reconstruction
- Enables study of brain-cell viability before anatomic changes occur in stroke and cancer treatment – within days!

Courtesy of Keith Thulborn and Ian Atkinson, Center for MR Research, University of Illinois at Chicago

©David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2009
University of Illinois, Urbana-Champaign

8

Advanced MRI Reconstruction

1 an Application Case Study

- magnetic resonance imaging
- **iterative reconstruction**

2 Acceleration on GPU

- determining the kernel parallelism structure
- loop splitting
- loop interchange
- using registers to reduce memory accesses
- chunking data to fit into constant memory
- using hardware trigonometry functions

a linear least squares problem

A quasi-Bayesian estimation problem:

$$\hat{\rho} = \arg \min_{\rho} \underbrace{\|\mathbf{F}\rho - \mathbf{d}\|_2^2}_{\text{data fidelity}} + \underbrace{\|\mathbf{W}\rho\|_2^2}_{\text{prior info}},$$

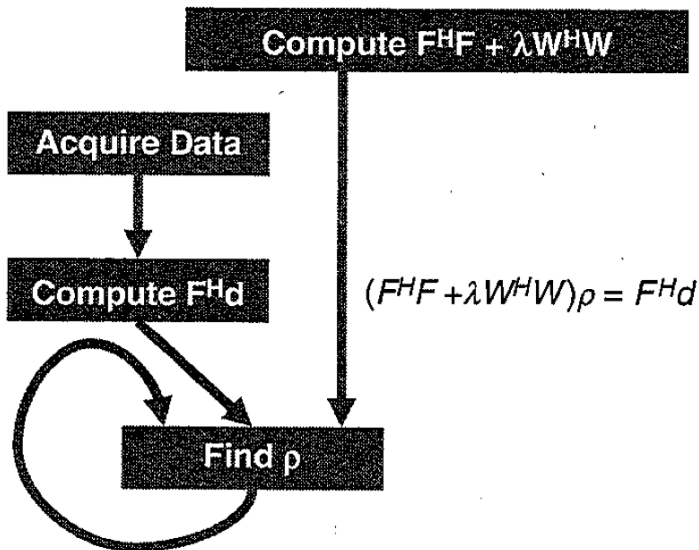
where

- $\hat{\rho}$ contains voxel values for reconstructed image,
- the matrix \mathbf{F} models the imaging process,
- \mathbf{d} is a vector of data samples, and
- the matrix \mathbf{W} incorporates prior information, derived from reference images.

The solution to this linear least squares problem is

$$\hat{\rho} = \left(\mathbf{F}^H \mathbf{F} + \mathbf{W}^H \mathbf{W} \right)^{-1} \mathbf{F}^H \mathbf{d}.$$

an iterative linear solver



three primary computations

The advanced reconstruction algorithm consists of

1
$$Q(\mathbf{x}_n) = \sum_{m=1}^M |\phi(\mathbf{k}_m)|^2 e^{i2\pi\mathbf{k}_m \cdot \mathbf{x}_n}$$

where $\phi(\cdot)$ is the Fourier transform of the voxel basis function.

2
$$\left[\mathbf{F}^H \mathbf{d} \right]_n = \sum_{m=1}^M \phi^*(\mathbf{k}_m) \mathbf{d}(\mathbf{k}_m) e^{i2\pi\mathbf{k}_m \cdot \mathbf{x}_n}$$

3 The conjugate gradient solver performs the matrix inversion to solve $(\mathbf{F}^H \mathbf{F} + \mathbf{W}^H \mathbf{W}) \rho = \mathbf{F}^H \mathbf{d}$.

The calculation for $\mathbf{F}^H \mathbf{d}$ is an excellent candidate for acceleration on the GPU because of its substantial data parallelism.

Advanced MRI Reconstruction

1 an Application Case Study

- magnetic resonance imaging
- iterative reconstruction

2 Acceleration on GPU

- determining the kernel parallelism structure
- loop splitting
- loop interchange
- using registers to reduce memory accesses
- chunking data to fit into constant memory
- using hardware trigonometry functions

computing $F^H d$

```
for(m = 0; m < M; m++)
{
    rMu[m] = rPhi[m]*rD[m] + iPhi[m]*iD[m];
    iMu[m] = rPhi[m]*iD[m] - iPhi[m]*rD[m];
    for(n = 0; n < N; n++)
    {
        expFHd = 2*PI*(kx[m]*x[n]
                      + ky[m]*y[n]
                      + kz[m]*z[n]);
        cArg = cos(expFHd);
        sArg = sin(expFHd);
        rFHd[n] += rMu[m]*cArg - iMu[m]*sArg;
        iFHd[n] += iMu[m]*cArg + rMu[m]*sArg;
    }
}
```

Consider the Compute to Global Memory Access (CGMA) ratio.

a first version of the kernel

```
__global__ void cmpFHD ( float* rPhi, iPhi, phiMag,
                        kx, ky, kz, x, y, z, rMu, iMu, int N)
{
    int m = blockIdx.x*FHD_THREADS_PER_BLOCK + threadIdx.x;

    rMu[m] = rPhi[m]*rD[m] + iPhi[m]*iD[m];
    iMu[m] = rPhi[m]*iD[m] - iPhi[m]*rD[m];

    for(n = 0; n < N; n++)
    {
        expFHD = 2*PI*(kx[m]*x[n] + ky[m]*y[n] + kz[m]*z[n]);
        carg = cos(expFHD); sArg = sin(expFHD);
        rFHD[n] += rMu[m]*cArg - iMu[m]*sArg;
        iFHD[n] += iMu[m]*cArg + rMu[m]*sArg;
    }
}
```

Advanced MRI Reconstruction

1 an Application Case Study

- magnetic resonance imaging
- iterative reconstruction

2 Acceleration on GPU

- determining the kernel parallelism structure
- **loop splitting**
- loop interchange
- using registers to reduce memory accesses
- chunking data to fit into constant memory
- using hardware trigonometry functions

splitting the outer loop

```
for(m = 0; m < M; m++)
{
    rMu[m] = rPhi[m]*rD[m] + iPhi[m]*iD[m];
    iMu[m] = rPhi[m]*iD[m] - iPhi[m]*rD[m];
}
for(m = 0; m < M; m++)
{
    for(n = 0; n < N; n++)
    {
        expFHd = 2*PI*(kx[m]*x[n]
                      + ky[m]*y[n]
                      + kz[m]*z[n]);
        cArg = cos(expFHd);
        sArg = sin(expFHd);
        rFHd[n] += rMu[m]*cArg - iMu[m]*sArg;
        iFHd[n] += iMu[m]*cArg + rMu[m]*sArg;
    }
}
```

a kernel for the first loop

We convert the first loop into a CUDA kernel:

```
__global__ void cmpMu ( float *rPhi,iPhi,rD,iD,rMu,iMu)
{
    int m = blockIdx * MU_THREADS_PER_BLOCK + threadIdx.x;

    rMu[m] = rPhi[m]*rD[m] + iPhi[m]*iD[m];
    iMu[m] = rPhi[m]*iD[m] - iPhi[m]*rD[m];
}
```

Because M can be very big, we will have many threads.

For example, if $M = 65,536$, with 512 threads per block, we have $65,536/512 = 128$ blocks.

a kernel for the second loop

```
__global__ void cmpFHd ( float* rPhi, iPhi, PhiMag,
                        kx, ky, kz, x, y, z, rMu, iMu, int N )
{
    int m = blockIdx.x*FHd_THREADS_PER_BLOCK + threadIdx.x;

    for(n = 0; n < N; n++)
    {
        float expFHd = 2*PI*(kx[m]*x[n]+ky[m]*y[n]
                           +kz[m]*z[n]);

        float cArg = cos(expFHd);
        float sArg = sin(expFHd);

        rFHd[n] += rMu[m]*cArg - iMu[m]*sArg;
        iFHd[n] += iMu[m]*cArg + rMu[m]*sArg;
    }
}
```

Advanced MRI Reconstruction

1 an Application Case Study

- magnetic resonance imaging
- iterative reconstruction

2 Acceleration on GPU

- determining the kernel parallelism structure
- loop splitting
- **loop interchange**
- using registers to reduce memory accesses
- chunking data to fit into constant memory
- using hardware trigonometry functions

loop interchange

To avoid conflicts between threads,
we interchange the inner and the outer loops:

```
for(m=0; m<M; m++)
{
    for(n=0; n<N; n++)
    {
        expFHd = 2*PI*(kx[m]*x[n]
                    +ky[m]*y[n]
                    +kz[m]*z[n]);
        cArg = cos(expFHd);
        sArg = sin(expFHd);
        rFHd[n] += rMu[m]*cArg
                  - iMu[m]*sArg;
        iFHd[n] += iMu[m]*cArg
                  + rMu[m]*sArg;
    }
}
```

```
for(n=0; n<N; n++)
{
    for(m=0; m<M; m++)
    {
        expFHd = 2*PI*(kx[m]*x[n]
                    +ky[m]*y[n]
                    +kz[m]*z[n]);
        cArg = cos(expFHd);
        sArg = sin(expFHd);
        rFHd[n] += rMu[m]*cArg
                  - iMu[m]*sArg;
        iFHd[n] += iMu[m]*cArg
                  + rMu[m]*sArg;
    }
}
```

In the new kernel, the n -th element will be computed by the n -th thread.

a new kernel

```
__global__ void cmpFHD ( float* rPhi, iPhi, phiMag,
                        kx, ky, kz, x, y, z, rMu, iMu, int M )
{
    int n = blockIdx.x*FHD_THREAD_PER_BLOCK + threadIdx.x;

    for(m = 0; m < M; m++)
    {
        float expFHD = 2*PI*(kx[m]*x[n]+ky[m]*y[n]
                           +kz[m]*z[n]);

        float cArg = cos(expFHD);
        float sArg = sin(expFHD);
        rFHD[n] += rMu[m]*cArg - iMu[m]*sArg;
        iFHD[n] += iMu[m]*cArg + rMu[m]*sArg;
    }
}
```

For a 128^3 image, there are $(2^7)^3 = 2,097,152$ threads.

For higher resolutions, e.g.: 512^3 , multiple kernels may be needed.

Advanced MRI Reconstruction

1 an Application Case Study

- magnetic resonance imaging
- iterative reconstruction

2 Acceleration on GPU

- determining the kernel parallelism structure
- loop splitting
- loop interchange
- **using registers to reduce memory accesses**
- chunking data to fit into constant memory
- using hardware trigonometry functions

using registers to reduce memory accesses

```
__global__ void cmpFHD ( float* rPhi, iPhi, phiMag,  
                        kx, ky, kz, x, y, z, rMu, iMu, int M )  
{  
    int n = blockIdx.x*FHD_THREAD_PER_BLOCK + threadIdx.x;  
    float xn = x[n]; float yn = y[n]; float zn = z[n];  
    float rFHdn = rFHD[n]; float iFHdn = iFHD[n];  
    for(m = 0; m < M; m++)  
    {  
        float expFHD = 2*PI*(kx[m]*xn+ky[m]*yn+kz[m]*zn);  
        float cArg = cos(expFHD);  
        float sArg = sin(expFHD);  
        rFHdn += rMu[m]*cArg - iMu[m]*sArg;  
        iFHdn += iMu[m]*cArg + rMu[m]*sArg;  
    }  
    rFHD[n] = rFHdn; iFHD[n] = iFHdn;  
}
```

Consider the improved Compute to Memory Access (CGMA) ratio.

Advanced MRI Reconstruction

1 an Application Case Study

- magnetic resonance imaging
- iterative reconstruction

2 Acceleration on GPU

- determining the kernel parallelism structure
- loop splitting
- loop interchange
- using registers to reduce memory accesses
- **chunking data to fit into constant memory**
- using hardware trigonometry functions

chunking k -space data into constant memory

Using constant memory we use cache more efficiently.

Limited in size to 64KB, we need to invoke the kernel multiple times.

```
__constant__ float kx[CHUNK_SZ], ky[CHUNK_SZ], kz[CHUNK_SZ];  
// code omitted ...  
for(i = 0; k < M/CHUNK_SZ; i++)  
{  
    cudaMemcpy(kx, &kx[i*CHUNK_SZ], 4*CHUNK_SZ,  
               cudaMemcpyHostToDevice);  
    cudaMemcpy(ky, &ky[i*CHUNK_SZ], 4*CHUNK_SZ,  
               cudaMemcpyHostToDevice);  
    cudaMemcpy(kz, &kz[i*CHUNK_SZ], 4*CHUNK_SZ,  
               cudaMemcpyHostToDevice);  
    // code omitted ...  
    cmpFHD<<<FHD_THREADS_PER_BLOCK,  
             N/FHD_THREADS_PER_BLOCK>>>  
        (rPhi, iPhi, phiMag, x, y, z, rMu, iMu, M);  
}
```

adjusting the memory layout

Due to size limitations of constant memory and cache, instead of storing the components of k -space data in three separate arrays, we use an array of structs:

```
struct kdata
{
    float x, float y, float z;
}
__constant struct kdata k[CHUNK_SZ];
```

and then in the kernel we use `k[m].x`, `k[m].y`, and `k[m].z`.

Advanced MRI Reconstruction

1 an Application Case Study

- magnetic resonance imaging
- iterative reconstruction

2 Acceleration on GPU

- determining the kernel parallelism structure
- loop splitting
- loop interchange
- using registers to reduce memory accesses
- chunking data to fit into constant memory
- using hardware trigonometry functions

using hardware trigonometry functions

Instead of `cos` and `sin` as implemented in software, the hardware versions `__cos` and `__sin` provide a much higher throughput.

The `__cos` and `__sin` are implemented as hardware instructions executed by the special function units.

We need to be careful about a loss of accuracy.

The validation involves a “perfect” image:

- a reverse process to generate “scanned” data;
- metrics: mean square error & signal-to-noise ratios.

The last stage is the experimental performance tuning.

references

This lecture is based on Chapter 8 (first edition; or Chapter 11 for the second edition) in the book of Kirk & Hwu.

- A. Lu, I.C. Atkinson, and K.R. Thulborn. **Sodium Magnetic Resonance Imaging and its Bioscale of Tissue Sodium Concentration.** *Encyclopedia of Magnetic Resonance*, John Wiley & Sons, 2010.
- S.S. Stone, J.P. Haldar, S.C. Tsao, W.-m.W. Hwu, B.P. Sutton, and Z.-P. Liang. **Accelerating advanced MRI reconstructions on GPUs.** *Journal of Parallel and Distributed Computing* 68(10): 1307–1318, 2008.
- The IMPATIENT MRI Toolset, open source software available at <http://impact.crhc.illinois.edu/mri.php>.

Versions of the F^H_d Algorithms on the GPU

- Base: pixel and scan data in global memory, accesses to scan data not coalesced, software sin/cos
- RegAlloc: pixel data in register file, scan data in global memory, accesses to scan data not coalesced, software sin/cos
- Coalesce: pixel data in register file, scan data in global memory, accesses to scan data coalesced, software sin/cos
- ConstMem: pixel data in register file, scan data in constant memory and constant cache, software sin/cos
- FastTrig: pixel data in register file, scan data in constant memory and constant cache, hardware sin/cos
- Tune: uses experimentally-tuned settings for loop unrolling, data tiling, and number of threads
- Multi: executes the tuned version on multiple GPUs

Performance of F^Hd Computation

